

NEBIUS GROUP NV NBIS

February 04, 2025 - 6:27pm EST

by

AAOI

			2025	2026
Price:	34.50	EPS	0	0
Shares Out. (in M):	235	P/E	0	0
Market Cap (in \$M):	8,100	P/FCF	0	0
Net Debt (in \$M):	-2,600	EBIT	0	0
TEV (in \$M):	5,500	TEV/EBIT	0	0

Description

Nebius Group (NBIS - NASDAQ)*AI Dark Horse Trading below Liquidation value***1) Introduction**

Over the past year, the AI boom has led to a shortage of high-power computing and data center resources, driving the rise of new business models like AI-focused cloud services. Thousands of AI startups and other enterprises are spending billions to develop AI solutions with hyperscalers and AI clouds as the foundational infrastructure providers. In the U.S., compute scarcity is so high even hyperscalers such as Microsoft are spending hundreds of millions (almost billions) per annum on these third-party AI clouds, while simultaneously investing in their own infrastructure. Additionally, Nvidia, the leading supplier of AI computing hardware (GPUs), is expanding its customer base and accelerating industry growth by offering preferential deals and capital to these smaller providers, such as CoreWeave, to help them increase their capacity while staying financially sustainable.

The advancements made by DeepSeek are deeply misunderstood (pun intended) by the market. Not only are the infrastructure costs likely hundreds of millions (possibly tens of thousands of H100s), but the ~\$5.5 million one-time training run cost cited by everyone isn't being properly compared for not only model style (distillation) but also the abilities (cited as a year or two behind US AI firms' distillation models). Today that same model could be run in the US by Anthropic for \$7-\$8 million. Nevertheless, the model is cheaper to run on lower tech hardware given its distillation, but this feature would likely expand AI use and not render many AI hardware and infrastructure players worthless.

Nebius Group (NBIS) is an emerging AI cloud provider (along with a few other business lines) supporting AI companies, with operations across Europe and North America. Originally, Nebius was part of Yandex, often referred to as the "Russian Google" due to its strong resemblance to Google in offering search, advertising, cloud, and various digital services within Russia and other formerly Soviet countries. After Russia's invasion of Ukraine in February 2022, Yandex faced significant challenges from the geopolitical fallout, which led to a halt in trading of its shares, its founder being placed on sanctions lists from Western governments and heightened regulatory pressure from Western nations and Russia. These issues forced Yandex to reevaluate its structure and geographic exposure (market and operational locales), particularly as its dependence on Western markets, partnerships, and investors was severely affected.

After an arduous two-year process, Yandex's founder was removed from sanctions lists in March and the Dutch parent company Yandex NV (YNV) sold its 28% stake in Yandex to a Russian consortium for \$5.4 billion in July, receiving \$2.8 billion in cash and 162.5 million YNV Class A shares. This move enabled Nebius to continue operating and growing in European markets while distancing itself from Yandex's Russia-based segment.

Now free of Russian ties, Nebius plans to expand beyond a collection of former Yandex assets and into a leading European AI cloud provider. As the company scales its AI infrastructure, it is well-positioned to leverage its experienced engineering team to develop a comprehensive full-stack AI solution, encompassing everything from compute infrastructure to deployment and monitoring. To address the DeepSeek concerns, this company is an AI services provider, closer to a AWS or MSFT Azure but built solely for AI development, it is not simply a GPU rental business like CoreWeave. This platform should provide a more defensible competitive advantage, moving beyond the market's concerns about earnings tied to the currently constrained GPU rental market, which could be commoditized in time.

With ~\$2.5 billion in cash (pro forma recent equity raise), a possible ~\$2.5 billion for the other divisions, no debt, and only a ~\$7.5 billion market cap, NBIS offers a remarkably cheap opportunity to invest in a successful founder and technology team at a critical inflection point in an industry with enough tailwinds to offset concerns on execution risk in their impressive growth plans. At current prices, even assuming a \$1 billion cash burn next year, Nebius's AI cloud business appears to be valued at ~\$2 billion - or ~2x ARR given the high end of management's "\$750-\$1,000 million ARR by year-end 2025" guidance, which we believe is highly achievable.

Given its non-core assets and flexible balance sheet, and its management's track record, Nebius has multiple opportunities (catalysts) to realize value and fuel growth plans in a manner accretive to shareholders. Additionally, the founder's close relationship with Jensen Huang seems to have secured Nebius greater access to affordable GPUs. This relationship also facilitated Nvidia's equity investment in Nebius during a late 2024 funding round, providing support similar to what CoreWeave has received.

2) Investment Thesis

We recommend investing in NBIS as the company executes its AI growth plans and achieves significant improvement in revenue and earnings, proving that it's more than a collection of leftover Yandex assets. Additionally, legacy concerns over Russian exposure and share ownership overhangs should fade representing another avenue towards material repricing of the core business.

At Nebius's current price of ~\$32/share, the market values its AI cloud business at ~2x ARR, even accounting for a considerable burn rate over the next year against ~\$2.5 billion in cash (pro forma recent equity raise), a minimal \$1 billion sale price for Avride (2x invested capital to date), and \$1.6 billion for its three other software businesses (Toloka, TripleTen, and a minority interest in ClickHouse). Importantly, Avride's expected \$150 million in necessary spending in 2025 is likely to be funded via a partnership investment, possibly Uber. And ClickHouse, while marked at a \$2 billion valuation in 2021, is one of the best performing assets of that vintage and is expected to undergo a funding round for over \$4 billion following the successful bid for a large-scale AWS contract. While we consider these two assets as freebies, the closest comps are Waymo and Snowflake, respectively.

Concerns over execution of these growth plans, cash burn from investment, and future levels of normalized earnings for AI cloud operations are reasonable but short-sighted, given management's track record, today's valuation, and the optionality embedded in Nebius's two main assets. Furthermore, the company has a well-defined roadmap to scale its AI business, modeled after CoreWeave,

which grew from zero to over \$2.3 billion in revenue in just a few years. CoreWeave's rapid success has earned it a premium valuation in private markets, with recent marks at a \$23 billion valuation (~\$33.5 billion EV) against ~\$2.3 billion in estimated 2024 revenue. The company is expected to IPO in 2025 with a \$35 billion valuation (\$45.5 billion EV). Based on a combination of 2024 figures and the potential valuation for the 2025 IPO, CoreWeave is valued around ~10x EV/Revenue.

If Nebius succeeds in becoming a leading AI cloud provider with a valuation on par with CoreWeave, an investor today gets that call option at an 80% discount to peers. If only Avride succeeds and is valued at one-third of Waymo's \$45 billion (with the AI cloud business worthless as GPU capacity goes into oversupply), Nebius still has a 2x option for which an investor pays next to nothing.

In fact, in a worst-case scenario, the liquidation value of Nebius's asset base is estimated at roughly \$31.25/share, roughly current price levels. Given the assets and optionality of the event path for its two main businesses, we believe there is asymmetry in the risk/reward ranging from today's price up to 685% upside, excluding Avride and ClickHouse.

There are several reasons for this current mispricing:

- a. **Investors old and new are absorbing information on an almost entirely new company, post the sale of Russian assets.** Yandex is a known entity to investors, the "Russian Google," even as it lives on in a much smaller, Russian-only capacity today. While the company was highly regarded and the biggest non-U.S. purchaser of Nvidia GPUs prior to 2022, Russian geopolitical risks kept many Western investors away. Nebius, essentially the remaining Western businesses of Yandex, is far less known despite many years of solid operations. Yet due to the recent resumption of trading after a two-year suspension, many former long-term Yandex shareholders are now faced with assessing the value of a set of businesses that differs significantly from what they initially invested in. On the other hand, the removal of Russian risks and exposure to two strong, growth-oriented businesses at critical inflection points is attracting new investors. In either case, long-term holders and new investors are in a period of investigation, determining whether the probability of success for the new businesses is appropriately priced in, or even worth taking a position on. This change in share ownership can create a mispricing and thus opportunity.
- b. **Entire company is a portfolio of call options, a setup typically mispriced by the market in the near term.** Nebius's significant investment in its AI cloud and Avride businesses effectively positions the company as a portfolio of real options, with both ventures representing new competitive plays built on repurposed legacy assets. The AI cloud business leverages Yandex's data center infrastructure in Finland, while Avride benefits from the strategic relocation of Russian engineering talent to Austin, Texas in early 2024. With both divisions currently operating at negative EBITDA and cash flow, traditional earnings-based valuations fail to capture Nebius's intrinsic value. As a result, the market overlooks these embedded growth options: A traditional earnings-based valuation is likely how former Yandex shareholders view the new company, leaving them inclined to exit. Additionally, substantial sunk costs effectively lower the exercise price of realizing future success, an advantage the market does not recognize, failing to apply an adequate premium to these growth options. As a result, Nebius's current valuation reflects a collection of out-of-the-money call options with significant upside potential and likely a far greater probability of paying off than the market price currently suggests.
- c. **Competitive advantages and earnings potential in the AI business misunderstood and overly discounted at current valuation.** At present, many view AI infrastructure as a GPU rental business with high capital intensity, rapid technological advancement/depreciation, little to no competitive differentiation, and therefore low ROIC. Even if these claims were accurate (which they largely are not) and GPU compute were to become commoditized, Nebius possesses several technological advantages that position it as a low-cost provider, rivaling the efficiency of the major hyperscalers. This business niche is clearly still in its early stages, and uncertainty around forecasted earnings is likely driving much of the potential mispricing. However, public investors can draw comparisons to CoreWeave, a business further along the development/growth curve, recently valued at over 10x its ~\$2.3 billion in 2024 revenue (up from ~\$500 million in 2023). In contrast, Nebius offers an opportunity to invest in a similar core business at just 2x ARR – an 80% discount based on the lower end of forward guidance. Moreover, management has additional strategic options that could unlock value, effectively setting a floor under the current share price regardless of future growth trajectories. As for the business itself, we believe that the company is going through two phases: (1) a high capital intensity "investment" phase of scaling up the AI infrastructure in a GPU supply-constrained environment, and then (2) a "moat-building" phase, in which it surrounds that infrastructure with comprehensive software tools to create a self-perpetuating Nebius AI ecosystem. Forcing a valuation hypothesis based solely on the current supply-demand imbalance in the AI compute market overlooks Nebius's robust infrastructure advantages and the considerable potential in its strategic vision.
- d. **Capacity growth embedded in guidance heavily discounted despite AI growth and subsequent compute scarcity.** AI clouds are becoming a critical addition to the data center market, driven by the distinct demands of AI workloads. Major hyperscalers are increasingly partnering with third-party providers to expand compute capacity, even as they invest heavily in their own infrastructure. For instance, Microsoft has committed \$10 billion to CoreWeave through 2030 to secure GPU availability, while Core Scientific reports tight GPU capacity through 2026, with customers funding capital expenditures upfront to accelerate buildouts. These dynamics lend credibility to Nebius's growth projections and its assertion that AI compute capacity will be quickly absorbed. To ensure near-term utilization, Nebius has aggressively priced its H200 GPUs at under \$2.50 per GPU/hour, compared to \$6-\$7 per GPU hour charged by peers. With this competitive pricing, Nebius's AI revenue has scaled rapidly, reaching ~\$200 million ARR in Q4 2024 – just one quarter behind capex deployment. The company plans to spend ~\$1 billion over the next year to more than quadruple GPU capacity, with \$400 million invested in Q4 2024 alone.

Also, Nebius is uniquely advantaged by brownfield (i.e. land that has already been developed and/or built on) expansion at its flagship data center in Finland, benefiting from Europe's greenfield (i.e. land that has not previously been built on or developed) development constraints. Those restrictions likely curtail broader AI data center growth across the region, further solidifying Nebius's market position and ability to fill its new capacity. Notably, access to Nvidia's next-generation Blackwell GPUs is critical to its 2025 growth plans, and Nebius's close ties with Nvidia CEO Jensen Huang position it as one of only three EU AI customers expected to receive Blackwells. In fact, Nebius's recent equity raise with Nvidia as a participant is a clear indicator that it's following the CoreWeave playbook. In the current supply-constrained environment, and given the linear relationship between capex and revenue, Nebius is on track to hit ~\$1 billion in ARR by year-end 2025. Any new evidence demonstrating its ability to secure this revenue could serve as a positive catalyst for the stock.

- e. **Concerns over cash burn rate and possible equity financing ignores massive cash balance and other creative options.** Nebius plans to invest approximately \$1 billion in GPUs and Avride development over the next year. Given high demand for GPU capacity, the company believes its AI growth is currently limited only by capital constraints. While Nebius's operations are currently cash-consuming, the company is on track to become EBITDA positive next year, driven by anticipated growth. With a cash balance of ~\$3 billion post-equity raise, Nebius is well-positioned to internally fund its growth plans for the coming year, while maintaining a healthy cash reserve. After 2025, the company aims to expand from ~60k GPUs to ~240k over the next few years, requiring approximately \$5 billion. Though ambitious, this target is not unreasonable; CoreWeave, for example, grew its revenue from \$500 million to over \$2 billion in a year, supported by a \$2.5 billion investment in 2023. Additionally, CoreWeave raised approximately \$9 billion this year, and now has a backlog of roughly \$17 billion through 2030, mostly backed by Microsoft. The company is simultaneously building out capacity in the highly coveted U.S. market, with ~35k GPU colocation in Kansas City. Nebius currently has no debt and could lever its balance sheet to finance its growth, if necessary. Again, CoreWeave shows the way – by issuing GPU asset-backed loans to fuel purchasing more GPUs, Nebius could bootstrap its way to 240k GPUs. Alternatively, the company could raise capital by selling Avride (valued at a minimum of \$1 billion) or other software businesses (conservatively valued at around \$500 million). With the founder holding approximately 13% of the company's stock, we anticipate a strong focus on the aforementioned options, with equity raises pursued only if they are expected to be accretive within a short timeframe.
- f. **Panic over DeepSeek appears to be based on incomplete data and lack of understanding of Nebius's business model.** DeepSeek, a Chinese AI startup, has recently unveiled its latest model, DeepSeek-R1, and a training cost of ~\$5.5 million. On the surface, this figure is substantially lower than the hundreds of millions typically associated with training large-scale AI models. However, this model appears to be distilled from LLaMA and Qwen and is not a "native" LLM built from scratch. DeepSeek learns from data created by other large language models and mimics their behavior to solve specific tasks. This allows it to achieve speed and efficiency in performing specific tasks while sacrificing the versatility and breadth of capabilities inherent to large base models. There are innovations in what they have done but it is evolutionary, not revolutionary. Additionally, the company is most likely downplaying their costs and hardware needs as access to tens of thousands of Nvidia GPUs would be problematic given the sanctions imposed. It is estimated that at the low end, a cluster would need to be tens of thousands and cost hundreds of millions to

stand up. Additionally, the model is more akin to models US competitors have had in market for over a year. If developed today, the US players could create their same models for only a few million more than DeepSeek, this is not a massive change in compute costs. Even if the concerns above were all valid, Nebius is an open-source AI cloud provider, and has already added DeepSeek on its platform. **While Nebius currently benefits from the compute supply/demand imbalance allowing them to scale GPU capacity quickly, it is positioning itself as an open-source provider to AI developers and is not a compute arbitrage company.**

- g. **Management underestimated, despite track record and substantive ownership.** Arkady Volozh co-founded Yandex in 1997, initially as a search engine. Over the next two decades, he expanded it into various sectors, including machine learning (now artificial intelligence) in 2009 and cloud storage in 2012. Yandex went public on the Nasdaq under the ticker YNDX in 2011 with an \$8 billion valuation, eventually growing to around \$30 billion by 2021. By then, Arkady's net worth exceeded \$2 billion, positioning him as the head of one of Europe's most prominent tech companies. However, in 2022, Yandex's market value declined significantly following the Russian invasion of Ukraine, leading to a suspension of trading in its shares for over two years. This past October, following a rebranding of its remaining assets as Nebius, trading resumed, with Arkady holding 90% of his net worth in NBIS, amounting to a ~13% ownership stake. Opportunities to invest alongside a founder of Arkady's caliber – who demonstrates exceptional personal commitment and a proven track record of leveraging his company as a platform for bold, innovative ventures – are rare and should not be overlooked by investors.

The largest micro risks to Nebius are execution issues with the data center expansion, the sales pipeline following said buildout, and methods of financing the growth that could be shareholder dilutive. The largest macro risks are the supply/demand dynamics for AI compute and utilization of said data center assets, the overall financial conditions of the market for raising capital, and technological changes that would reduce the value of a GPU-centric industry.

3) Business Analysis

Overview: Assets of the former “Russian Google” coalescing into a formidable AI platform

Yandex was founded in 1997 by Arkady Volozh and Ilya Segalovich, both of whom had backgrounds in technology and computer science. The company started as a search engine project named Arkadia, but this soon evolved into Yandex (which stood for “Yet Another iNDEXer”). The search engine was tailored to the Russian language and its unique syntax, which was not being adequately served by Western search engines like Google or AltaVista at the time.

Over the next 20 years, Yandex expanded into services such as maps, email, and online advertising. Notably, the company started a machine learning operation in 2009 and a cloud storage service in 2012. Yandex Self-Driving Group (SDG) was established in 2016, moving to Austin, TX as Avride in 2024. Yandex Cloud was officially launched in 2018, with its Finnish data center being the company's prime asset outside of Russia.

Following the Russian invasion of Ukraine in 2022, trading in Yandex was suspended, the founders were subject to sanctions, and employees left Russia in droves. During the next two years, the business's Russian assets were split off from its Western ones. The Russian assets, Yandex proper, were sold to a Russian consortium for \$5.4 billion in mid-2024 as described earlier. Shares in this Western asset group resumed trading in late 2024 under the name Nebius Group NV (NBIS). It consisted of the AI cloud (Nebius), three software assets (Toloka, TripleTen, ClickHouse) and an autonomous driving unit (Avride).

Nebius AI Cloud - Compute and Services

Nebius AI is currently a “compute-as-a-service” or AI infrastructure business. The company designs its own servers (outside of Nvidia GPUs) and AI cluster layout in a data center and then rents GPU capacity to companies developing AI programs and services. Currently, the company has over 40 customers, many of which are AI-native.

The company's main compute asset is the former Yandex Finnish data center which is being repurposed for AI use cases. Nebius plans to expand capacity in that data center by 3x to ~75 MW, or 60k GPUs, by early 2026. Additionally, the company is running a co-location in Paris with Equinix, and another in Kansas City with Patmos. The Kansas City cluster should start at 5 MW early next year, and then scale to 40 MW, or 35k GPUs, over the next year or two. The Equinix cluster in Paris is capable of ~30 MW, or 25k GPUs, and should be the first in Europe to receive Blackwells.

At present, NBIS is running ~20k GPUs with plans to reach ~60k in 2025 and ~240k in 2027+. The CEO has hinted that is may be able to build “hundreds of MWs in 2025” and “reach 1 GW in 2026”, which would likely be a massive pull forward of capacity backed by banks and other financial players.

AI Services

Nebius offers not only AI cloud services but is building an open-source AI platform and application layer. These are scalable but in early stages, shown below in purple.




Figure 1: Nebius October IR Presentation.

Additionally, Nebius is building Nebius Studio, an API SaaS for open-source models, shown below:




Figure 2: Nebius October IR Presentation.

These developing products and services are tailored for the inference market.

Competitive Advantages

AI Infrastructure

Nebius manages the entire AI infrastructure stack, including in-house server design, data center operations, and a proprietary cloud platform. With highly specialized hardware and designs, this approach allows them to optimize performance, which then leads to lower costs. Its team of seasoned system architects, electronic engineers, developers, and mechanical and thermal engineers were responsible for creating Yandex's 200 MW of cloud capacity in the past. The Nebius team has over 850 AI/ML/cloud engineers, most of which are in the Netherlands or Israel.

The company has implemented specialized cooling and heat co-generation systems at its Finnish data center, capturing heat generated by AI servers and repurposing it for energy use. This heat can be converted into power via mechanical processes (e.g. steam turbines or heat pumps) or directly utilized, such as heating water for nearby homes, shown below:




Figure 3: Data Center Frontier.

By reusing heat (energy) from AI operations either recirculated inside the data center or in external applications, Nebius achieves exceptional efficiency. Up to 75% of the heat generated from its AI operations can be recovered. This is reflected in its Power Usage Effectiveness (PUE), which rivals that of top hyperscalers, as shown below:




Figure 4: Nebius October IR Presentation.

Importantly, Nebius-built data centers will have performance enhancements from their own designs as well as this cooling/co-gen capability, meaning their PUE is replicable and not just a function of low ambient temperatures due to their northern location in Finland. Overall, Nebius management believes its solutions result in a 20%-25% lower total cost of ownership compared to that of other GPU providers. **If compute does get commoditized, the lowest cost provider wins, and here we see Nebius alongside GCP and MS Azure.**

For co-location data centers, Nebius is implementing its custom designs but not the cooling/co-gen, so the PUE for these locations will be a higher, but still impressive, 1.13 PUE, nowhere near the ~1.5 global average.

Nvidia Relationship – GPU supply, advantageous pricing, possible investment to fuel growth

Nebius, in its prior form as Yandex, was one of Nvidia's largest non-US customers. The company has already indicated it is not only going to be one of the first providers of Blackwells in Europe but also that it can purchase them at very competitive prices compared to peers. In late 2024, NBIS announced pre-sales for 22k Blackwell GPUs for use in early 2025. Considering the company stated it would spend \$400 million in 4Q for GPUs, we can assume it was for these Blackwells. If such is the case, the per GPU cost would come to less than \$20k per GPU, a significant discount to peers. We think this relationship is sustainable and should help jump start Nebius's AI growth. To that point, as of late 2024, the company raised \$700 million with Nvidia participating. More on this below.

Nebius AI's Strategy – CoreWeave Growth Plan Followed by Full Stack Open-Source Platform

CoreWeave started in 2017 in Roseland, New Jersey. It quickly grew from zero revenue to ~\$2.3 billion in 2024 by targeting high-demand, GPU-intensive markets and leveraging a close partnership with NVIDIA. Initially, CoreWeave focused on GPU computing for cryptocurrency mining, but it pivoted to cloud-based AI infrastructure as the AI boom created massive demand for GPUs.

The company financed its growth through a combination of private investments, debt financing, and strategic customer agreements. On debt financing, CoreWeave used its GPUs as collateral to receive a better interest rate than it could have achieved using only its then-current financials, which were not generating cash flow during the company's hyper-growth phase.

Nvidia played a crucial role in CoreWeave's success by providing early access to cutting-edge GPUs, preferential pricing, and direct financial support through an equity stake. This relationship also gave CoreWeave credibility, helping it attract customers and investors while ensuring priority access to GPUs. Its deep integration with Nvidia's hardware and software ecosystems ensured it remained at the forefront of AI advancements, reinforcing its market position. It also gives Nvidia a strong customer outside of the hyperscalers.

Notably, Nebius may be mispriced due to a small group of investors mistakenly perceiving this relationship as a nefarious effort by Nvidia to inflate an AI bubble. This misperception similarly affected CoreWeave last year, leading to unwarranted skepticism about its partnership with Nvidia.

With those pieces in place, CoreWeave was able to attract Microsoft as a large customer, generating \$10 billion in revenue over the next six years. Supposedly, the company is already profitable and ramping to \$8 billion in revenue next year, with a 10% minority stake floated via IPO planned.

Nebius's future growth is likely to follow a two-phase strategy, with the first in the same vein as CoreWeave. In the initial phase, the company will likely raise capital through debt financing or by divesting non-core assets to fund GPU capacity expansions. Any equity investments will likely be backed by strategic partners. In the AI infrastructure business, capital intensity acts as the primary barrier to entry, and the company has multiple levers to clear this hurdle. Next, Nebius's specialized servers and workflows for AI provide differentiated value to customers at a lower cost, giving it a competitive edge once scaled. Its close partnership with Nvidia could further fuel growth, supported by Nvidia's late 2024 equity investment, as well as continued preferential GPU pricing. Within a year or two, Nebius could potentially scale enough to reach a market position comparable to CoreWeave's. However, unlike CoreWeave, Nebius's seasoned engineering team is intentionally building for a second, even more ambitious phase.

In this second phase, Nebius aims to establish an open-source software ecosystem centered around its large-scale GPU infrastructure, offering cost-effective, full-stack AI solutions. This strategy would not only more effectively monetize older node GPUs but also protect the company against competition based solely on access to the best GPUs and their supply and demand. By integrating tailored AI solutions

with hyperscale-level compute capabilities, Nebius can differentiate itself from hyperscalers, which often lack customized offerings. This approach also positions Nebius to outcompete the fragmented market of point solutions and secure long-term leadership in AI infrastructure and software.

Full-Stack AI Ecosystem - What Does It Look Like, Why Does It Matter?

If Nebius evolves to a full AI stack with a tokenized pricing model, it would fundamentally change its business as an AI cloud service. The full stack is shown below. At present, Nebius is in layer 1 with its infrastructure and partly in layer 2 with Toloka:




Figure 5: Menlo Ventures.

With infrastructure and then services and software (either built internally by Nebius or through an API), a potential AI user would have a one-stop shop to build/use an AI, with pricing based on tokens. An "AI token" refers to the smallest unit of data that an artificial intelligence model processes, particularly in natural language processing (NLP), where text is broken down into individual pieces (such as words or parts of words) for the AI to understand and base its responses on. The tokenization model simplifies billing, allowing customers to pay only for the exact resources they use. Pricing is based on how many units of words or numbers one wants to analyze (e.g. 1,000 tokens per \$1). A full-stack AI service would simply price its tokens above the cost to run a token, playing into Nebius's low-cost capabilities.

Tokens also provide greater flexibility and scalability, enabling customers to allocate resources as needed. For example, customers could increase their token usage during training phases and scale it back during maintenance. This would especially benefit startups and research institutions with dynamic AI workloads, driving significant cost savings by avoiding over-provisioning.

Integrating tokens into an expanded ecosystem - including elements such as AI tools, data labeling services, and software products - enhances its value proposition and deepens customer engagement, creating a wider moat than Nebius would have from infrastructure alone. Importantly, the economics of this ecosystem would disassociate AI customers from the underlying compute hardware, meaning customers would pay per unit of what they would like to compute versus paying specifically for the compute hardware, letting Nebius distance itself from the GPU-only infrastructure treadmill model.

Toloka

Toloka specializes in creating high-quality labeled datasets for AI and machine learning by offering a range of services tailored to various stages of model development. Its core capabilities include data labeling for images, text, video, and audio, enabling supervised learning tasks such as object detection and sentiment analysis. The platform supports AI fine-tuning and evaluation by providing datasets for Reinforcement Learning from Human Feedback (RLHF) and fine-tuning large language models (LLMs) through preference collection, ranking, and model comparison.

Toloka also integrates human input into AI workflows with its "human-in-the-loop" approach, ensuring accuracy and relevance. The division's AI clients include Amazon, Microsoft, ServiceNow, Hugging Face, and others:




Figure 6: Nebius October IR Presentation.

TripleTen

TripleTen is a leading edtech platform that specializes in reskilling and upskilling individuals for successful careers in technology. The company offers immersive, project-based learning experiences tailored to equip students with in-demand skills in software engineering, data science, cybersecurity, and other technology-driven fields.

Avride

Avride is an autonomous driving technology company focused on developing solutions for self-driving cars and delivery robots. The company leverages shared technologies, including advanced machine learning, computer vision, and sensor integration, to create scalable and efficient systems for various applications. Avride's autonomous vehicles are designed for ride-hailing services, food and grocery delivery, and logistics, while its delivery robots cater to last-mile deliveries, such as restaurant orders and small e-commerce packages. The company has achieved significant milestones, including over 47k passenger rides and 200k robot deliveries, with operations in diverse environments including the U.S., South Korea, and Europe.




Figure 7: Nebius October IR Presentation.

Avride's strategic partnerships, such as with Uber for autonomous delivery and mobility, position it to scale rapidly in high-growth markets. By integrating autonomous technology into production vehicles and leveraging efficient design and operations, Avride aims to lead in the autonomous vehicle and delivery sectors, offering innovative and cost-effective solutions.

Notably, Avride has completed one-third the kilometers of testing and service compared to Waymo, a key factor underlying our rule-of-thumb valuation for Avride relative to Waymo, as outlined in the thesis. **Notably, the CEO has stated that the \$150 million in funding needs for Avride over the next year or two is "taken care of", meaning a strategic partner, we think Uber, should come in with capital this year.**

ClickHouse

ClickHouse is an open-source, columnar database management system designed for high-performance analytics and real-time query processing. Originally developed by Yandex in 2009, open-sourced in 2016, and spun off in 2021, ClickHouse is optimized for OLAP (OnLine Analytical Processing) workloads, making it ideal for scenarios requiring fast data aggregation and analysis over large datasets.

Nebius has a 28% ownership stake in ClickHouse, and the latest available valuation of the company was \$2 billion in 2021. Tech asset valuations peaked in 2021, and many think the asset isn't worth much. **Doing even a modest amount of due diligence on this company, with the new AWS contract in hand, ClickHouse is the best performing investment from the 2021 vintage of investments from Coatue etc. The business is likely to go through a funding round this year post the AWS contract at a valuation above \$4 billion, putting Nebius's stake at \$1+ billion.**

Management: Impressive founder CEO "doing it all over again" after sanctions and considerable turmoil; Owns 13% of NBIS, roughly 90% of his net worth.

Under Volozh's leadership, Yandex grew into Russia's leading search engine and one of Europe's largest internet companies. The company expanded its services to include e-commerce, transportation, and other technology-driven sectors.

In 2014, following the Russian annexation of Crimea, Volozh relocated to Tel Aviv, Israel, with his family. He continued to lead Yandex until 2022, when he stepped down as CEO after being sanctioned by the European Union over Russia's invasion of Ukraine. Volozh holds Israeli citizenship and has been residing in Israel since 2014. He has publicly condemned Russia's invasion of Ukraine, a stance that facilitated his removal from EU sanctions lists in early 2024.

Volozh founded Nebius Group in 2024, an AI-focused company headquartered in Amsterdam and containing the non-Russian assets spun off from Yandex. Yandex proper was sold to a Russian consortium for ~\$5.4 billion, comprising \$2.3 billion in cash and 162.5 million shares (which are held in treasury against ~200 million shares outstanding). Volozh holds ~13% of NBIS equity and has stated it is 90% of his net worth. With the voting rights between the ~163 million Class A and ~35.5 million Class B shares, Volozh has ~60% voting power.

Notably, we believe it's more likely the company will retire most of the 162.5 million shares received from its Yandex sale rather than sell them to raise cash, with the remainder used as employee stock compensation that would reduce its current \$400 million Dutch tax bill.

Customer Dynamics: AI compute still in high demand, full-stack open-source AI platform would create value beyond GPU rentals

Nebius currently serves 40 customers providing AI Infrastructure-as-a-Service (IaaS). Customers must manage their own AI stack, from training models to deploying them. Most of NBIS's customers are AI-centric, working on training and fine-tuning large language models, multimodal systems, and domain-specific AI solutions. The company is pricing its GPUs 50% below those of peers to be flexible and cost-efficient, ensuring capacity utilization. One of the company's customers is Mistral AI, one of the largest AI firms in Europe.

Nebius is also targeting an expanded customer base beyond IaaS. Once it scales its infrastructure and builds out its software offerings to be a full-stack AI platform, it plans to target GenAI app builders and non-AI-centric enterprises, shown below:




Figure 8: Nebius October IR Presentation.

Supplier Dynamics: Close relationship with Nvidia founder an avenue for plenty of GPUs at structurally lower prices; tenured cloud and machine learning (AI) engineers in large numbers from legacy Yandex

Compute

Multiple sources expect AI GPU demand to double from 2024 to 2026. Mismatches in supply and demand should continue. Nebius's close relationship with Nvidia, and that company's desire to support competitors to hyperscalers, should aid in GPU supply and pricing. Current estimates on Blackwell prices for Nebius range from \$25,000 to \$30,000 per GPU. To scale from 60,000 GPUs in 2025 to a total of 240,000 GPUs, Nebius will likely need approximately \$5 billion in funding over the next few years.

Recent estimates, as noted earlier, suggest that Nebius may be paying less than \$20,000 per GPU, which would reduce our projected \$5 billion capital outlay.

Infrastructure - Data Centers and Power

Based on industry sources, we believe CoreWeave is paying \$7-\$10 million per MW for brownfield datacenter capex in the U.S. Additionally, greenfield expansion is estimated at \$10-\$15 million per MW. In Europe, the cost per MW is estimated to be ~20% higher.

One of the largest impediments to growth is power supply. With larger data centers and a bigger drain on local resources, Nebius may engage in similar deals with power providers in Europe and U.S as the leading hyperscalers are doing today.

Finland is a desirable location for data centers, given the developable land, available power and workforce. Google's main European data center is located in Finland.

Engineers

Nebius has 850 engineers with over ten years of cloud, AI, ML, and LLM experience. Additionally, it has no employees in Russia, most being in the Netherlands, Serbia, Israel, and the U.S.

Competitor Dynamics: Low power usage effectiveness (PUE) on par with hyperscalers is a huge differentiator for AI workloads

Hyperscalers like AWS, Azure, and Google Cloud offer extensive global scale and integration, making them appealing for enterprises already embedded in their ecosystems. These providers also benefit from established customer bases and broader service portfolios, which provide convenience and compatibility. As they can buy GPUs by the tens of thousands, they can exercise buying power better than smaller players, and/or design their own chips to create a substitute.

AI cloud providers such as CoreWeave and Lambda are strategically addressing hyperscalers' demands by offering surplus GPU capacity while catering to AI-focused customers. Nebius stands out with its development of specialized software solutions that add a competitive layer to its infrastructure stack. Additionally, these emerging players are well-positioned to meet local regulatory and privacy requirements, particularly in regions like Europe and Asia. By prioritizing cost-effective solutions and ensuring reliable GPU availability, they appeal strongly to budget-conscious customers seeking value and dependability.

As mentioned above, Nebius differentiates itself with its industry-leading Power Usage Effectiveness (PUE) of 1.13, positioning it among the most energy-efficient operators. This efficiency directly contributes to cost savings, sustainability, and competitive pricing in the GPU rental market.

Market Trends: AI Experiencing rapid growth with large potential market

AI

The GPU-as-a-Service and AI Cloud market is expected to expand from ~\$33 billion in 2023 to ~\$260 billion by 2030 (for a ~35% CAGR), according to the Nebius internal estimates:




Figure 9: Nebius October IR Presentation.

Last year, Nvidia's VP of Enterprise Computing stated management believes the market for AI could reach over \$600 billion, half on chips and systems, and the other half for AI software and services. Many other experts have cited the large potential for AI; there's not much point to listing all their sky-high estimates here. The market is very large, but as it is still in the early stages, it remains somewhat amorphous. This quality should let Nebius build a sizable position over time.

Training vs. Inference

AI training involves teaching models to recognize patterns by processing vast datasets. This phase is computationally intensive, requiring significant resources to develop accurate models, and with customers requiring the latest and greatest GPUs. Training has been the primary driver of AI infrastructure spending to date. However, the market should shift to more inference as trained models get deployed.

AI inference is the application of trained models to new data, enabling tasks like image recognition or language translation. Inference operations are generally less resource-intensive than training, but occur more frequently across various devices and applications. As such, inference markets would allow fully-depreciated GPUs to still "earn a living" after they are no longer leading-edge chips used in training.




Figure 10: Nebius October IR Presentation.

Nebius plans to focus its efforts on the inference market, as in time it should be larger than the training market and less resource-intensive.

4) Why now?

Nebius is at a key inflection point with a news event that has introduced considerable confusion: It's currently undervalued compared to its peers, while purposely following those same peers' roadmap to scale its AI business. The name is early on its discovery process with investors, with no sell-side coverage yet, and it has multiple catalysts in the next 12 months:

- Revenue growth should increase in lockstep with 4x-5x GPU capacity growth through 2025, and any new information that points to this growth should be a catalyst.
 - AI ARR from ~\$200 million to \$750-\$1,000 million by year end-2025 highly likely given GPU supply deficit.
 - Upside possible to 1GW or ~\$10 billion in ARR based on the CEO's assertions. Dare we say they up the ante and guide to 1.21 GIGAWATTS!!!! Arkozy guides to flux capacitor like magic here: <https://t.co/3iyZxluq2g>; note this is ~4x prior "medium-term" guidance by next year! Make of that what you will.**
 - Nebius should announce an anchor customer in 2025. **Notably, NBIS's Finland data center is ~100km from Google's.**
- Multiple options for value-accretive growth financing should also act as catalysts.
 - Nvidia-backed equity raise already done. PIPE of \$700 million done at \$21/share, a premium to the previous share price, shows enthusiasm for the company.
 - Expect GPU-backed loans to bootstrap GPU capacity increases, similar to CoreWeave. Or a bank backstopped facility like Applied Digital's deal with Macquarie.
- Non-core asset should be monetized or given new valuation marks over the medium term.
 - Avride (minimum \$1 billion): Mid to late 2025, with \$150 million investment.
 - ClickHouse: New mark to market valuation of \$4 billion, with growth asset could be worth \$8 billion. Nebius's stake would be marked at ~\$1.1 billion or ~\$2.2 billion in each case.
 - TripleTen (\$100-\$200 million): 2x-4x 2025's \$50 million ARR.
- CoreWeave IPO in 2025 could price at 6x-10x EV/Revenue, serving as a catalyst for NBIS as well.

Nebius currently trades at ~2x YE'25 ARR at an 80% discount to its closest peer, with a large increase in growth coming, multiple levers at

the asset level and balance sheet to accelerate growth, and backed by an executive team second to none.

Looking ahead, Nebius has the potential to evolve into a full-stack AI platform, positioning itself as an AI-focused hyperscaler with profit margins comparable to those of industry leaders like AWS, GCP, and Azure. Even modest valuations (EV/EBITDA multiples in the teens) on Nebius's plans for 240k GPUs with GPU per hour rates well below market into perpetuity, yield per-share values of 2x the current price. More details on this in the Valuation section.

-

5) Risks/Thesis Pressure Points

GPU Supply/Demand Risk. At present, GPU availability is low, and providers of GPUs like AI clouds can charge prices well in excess of their cost. With current useful lives of ~4 years, current prices allow for payback within a year or so, making this environment highly economic. Should AI workloads require less compute in the future or experience a downturn in overall levels of utilization, this supply imbalance could flip and prices of leading edge GPUs would drop.

Execution Risk. Nebius has to first purchase GPUs, integrate them into a data center, and sell the capacity quickly. Much of this thesis is predicated on revenue growth (as well as subsequent margin expansion) that requires smooth execution by multiple teams.

Technology Risk. When new GPUs come into the market or older ones become more widely available, the rental prices of the older ones drop. For instance, H100 rental prices experienced a 25% decrease as H200's came onto the market. The current pace of development has new GPUs hitting the market every 1-2 years. However, useful lives of GPUs are expected to increase with full liquid cooling, which has led to an 80% reduction in GPU failure rates.

Financing Risk. Nebius requires liquidity to fund its ambitious growth plans, through equity, debt, and asset sales. Should the markets shun AI or experience a tightening of liquidity, the company's growth plans would either become more expensive to finance or severely diminish.

6) Valuation

Nebius is scaling its AI infrastructure, and over time should provide full-stack open-source AI software with a token pricing model. We assume the business at a minimum should look similar to AWS and Azure at scale. We give no credit to Avride or even AI software services that could be layered on top. The other software businesses assume no growth.

The following are our Base Case assumptions for the company:

1. **Company Guidance:**
 - a. AI ARR: \$750 million to \$1 billion by YE'25, GPU capacity increases 4x-6x.
 - b. 60k Blackwells, scale to 240k, capital availability dependent.
 - c. 2025 Capital expenditures: AI: \$600 million to \$1.5 billion, Avride: \$150 million opex/capex
 - d. Adjusted EBITDA positive in 2025 but below "break-even" margins.
 - e. Other Revenues: Toloka: \$50-\$70 million, TripleTen: \$40-\$60 million.
 - f. \$400 million Dutch tax bill largely mitigated via share-based compensation issuance from ~162.5 million shares held in treasury from Yandex sale.
2. **Long-Term Internal Assumptions:**
 - a. Scale to 240k GPUs over the next 5 years: \$5 billion capital cost.
 - b. Revenue growth: \$700 million in 2025 (~\$1 billion ARR by YE'25) to \$6 billion by 2030.
 - c. Leading edge GPU pricing: \$4 per hour, 30%-40% below peers.
 - d. Trailing edge GPU pricing: current prices decline 20% annually after 3 years.
 - e. GPU useful life: ~70% utilization, 5 years, assumes benefit from 100% liquid cooling.
 - f. Gross Margins: 65%, below CoreWeave's estimated 85% and cloud providers 70%+.
 - g. EBITDA margins: scale from 2% in 2025 to a normalized 35% by 2028, roughly on par with AWS and below Azure.
 - h. Normalized capex in addition to GPU spending: ~8% of revenue.
3. ~15x normalized EV/EBITDA multiple (equates to ~5.75x EV/Revenue and ~20x EV/FCF). Our reasoning behind these multiples is in the Peer Analysis section below.
4. Discount rate of ~10% (large cap).

The following are our Bull and Bear Case assumptions for the company:

1. **Bull Case Internal Assumptions:**
 - a. Scale to 1,000 GPUs over the next 3 years - \$15 billion capital cost. **(1.21 GIGAWATTS! Great Scott!!)**
 - b. Revenue growth: \$700 million in 2025 to \$20+ billion by 2030.
 - c. Leading edge GPU pricing: \$4 per hour.
 - d. Gross Margins: 70%.
 - e. Normalized EBITDA margins: 40%-45% on par with Azure's Intelligent Cloud division.
 - f. Normalized EV/EBITDA multiple: 20x.
 - g. Bull Case valuation implies 20% share of \$260 billion AI market, with mid-teens FCF margins at 17.5x EV/FCF.
2. **Bear Case Internal Assumptions:**
 - a. Liquidate company as is.
 - b. \$2.6 billion in cash (\$3 billion pro forma equity raise less \$400 million 4Q capex)
 - c. Finnish Data Center - \$2.1 billion, 65 MW at ~\$30K per MW.
 - d. Avride: ~\$1 billion (2x invested capital).
 - e. Other Software: \$1,600 million (\$1,100 million for ClickHouse, \$300 million for TripleTen, \$200 million for Toloka).

A summary of the three valuation scenarios is below:

Valuation Ranges - Ex. Avride			
Scenario	Bull	Base	Bear
Value/share	\$249.50	\$60.50	\$31.25
Upside/(Downside)	685%	90%	(2%)
Per annum IRR	51%	14%	(0%)

To clarify, our Base Case valuation excludes any potential upside from increased contributions of AI software solutions and assigns no value to Avride, whether through a sale or future financial performance. In contrast, our Bull Case accounts for the 1GW potential cited by the CEO recently and some uplift in margins from AI software but still does not attribute any value to Avride or ClickHouse.

Assigning probabilities of 40%, 40%, and 20% to the respective scenarios above results in an expected average target price of approximately **\$130/share**, equating to an annualized **IRR of ~35%** over a five-year timeframe.

If one were to include Avride, we believe the uplift to per share value could look like this:

Valuation Ranges - With Avride and ClickHouse			
Scenario	Bull	Base	Bear
Avride/CH Value (\$B)	\$17	\$3	\$1
Value/share	\$322.75	\$73.75	\$31.25
Upside/(Downside)	915%	132%	(2%)
Per annum IRR	59%	18%	(0%)

Peer Analysis

Nebius has no direct public comparables yet, which is one reason why the stock is mispriced. The cloud divisions of Amazon, Google, and Microsoft are close competitors in a sense. Private companies CoreWeave and Lambda Labs are direct competitors. To triangulate a fair valuation, we used an EV/EBITDA framework based on the public cloud companies, while citing upside based on CoreWeave's private valuations and possible IPO next year. EBITDA is a more relevant metric for assessing operational performance, as it accounts for the significant differences in capital structure and life-cycle stages between Nebius and its peers.

Trading Comparables

EV/EBITDA

1. Amazon (AMZN): ~15x normalized EV/EBITDA; currently at ~14.25x FTM.
2. Google (GOOG): ~12x normalized EV/EBITDA; currently at ~12x FTM.
3. Microsoft (MSFT): ~20x normalized EV/EBITDA; currently at ~20x FTM.

While not a perfect fit, the normalized EV/EBITDA multiple comes in at ~15x. Considering NBIS isn't generating positive EBITDA at the moment, this normalized multiple is based on our assumption the core business could look similar to those peers' cloud divisions.

EBITDA Margins

1. Amazon (AMZN): 25%-30% EBIT, ~8.5% depreciation estimate - EBITDA 35%+.
2. Microsoft (MSFT): 40%-45% EBIT, ~10% depreciation estimate - EBITDA 50%+.

These margins are a best-efforts estimate for a scaled cloud division. Operating income is easily observed, but D&A is an estimate. For our Nebius valuation, we assume the low end.

Private Company Analysis - CoreWeave

CoreWeave's financial information and private valuations:

- a. Revenue: 2024: ~\$2.3 billion, 2025: ~\$8 billion. Backlog: \$17 billion.
- b. GPU pricing \$6-\$7 per hour.
- c. Gross margins: 85%.
- d. ~\$10.5 billion in debt raised.
- e. Possibly ~270 MW of capacity today and plans to reach 500 MW.
- f. Mark-to-Market valuation: 2024: \$23 billion (EV ~\$33.5 billion) - 14.5x EV/Revenue
- g. Estimated IPO: \$35 billion (EV ~45.5 billion) - ~5.7x EV/Revenue

We could put NBIS at the midpoint of these EV/Revenue valuations - a multiple of approximately 10x. However, based on CoreWeave's EV/MW valuation (270 MW valued at ~\$150 million per MW), NBIS, with 240 MW at scale, would be valued at approximately \$145/share, slightly below our Bull Case figure.

If we assume CoreWeave's valuations price in growth to its full 500 MW capacity, or ~\$80 million per MW, then NBIS with 240 MW at scale would be worth ~\$80/share, slightly above our Base Case figure.

7) Market Expectations/Perceptions

Nebius currently lacks coverage from sell-side analysts, but this may soon change given its close relationship with Goldman Sachs. While special situation buy-side analysts are actively tracking the company, it remains relatively under the radar for most investors due to its status as a "new company." Most legacy Yandex holders, who may have been stuck in stock, have probably exited by now. Current confusion as to the outcomes in the industry following the DeepSeek advancements present an opportunity.

In a market that places a premium on AI growth, Nebius offers a potentially undervalued opportunity to invest in this trend. Its founder, Arkady Volozh, is highly regarded for his expertise, adding a "back the jockey" dynamic that bolsters confidence in the company's prospects.

8) Downside Protection - Where's the Margin of Safety?

The company's downside is well protected. A liquidation analysis in our Bear Case puts a floor in the lower \$30s for the shares, which is a little above the current price.

Avride requires \$150 million over the next year, and the company has indicated that it has a strategic to invest alongside this year. We pencil in \$1 billion as a possible sale price, but do not include it in our valuation. Considering its competitiveness with Waymo, once scaled it could be worth far more than \$1 billion. Same with ClickHouse, a new valuation mark of \$4 billion would have Nebius's stake at over \$1 billion.

9) Conclusion

Nebius represents a compelling investment opportunity at a critical juncture in the AI infrastructure market, with multiple call options on what could be a leading full-stack open-source AI platform. Still seen as a newcomer, the company has extensive experience in AI cloud services. With a highly capable management team led by founder Arkady Volozh and a proven ability to execute complex strategic pivots, Nebius is well-equipped to scale rapidly while navigating the complexities of a fast-evolving market.

Despite its origins in Yandex and the lingering perception of its Russian ties, Nebius has successfully redefined itself as a Western-focused, growth-oriented enterprise, operating independently with a significant cash balance and no debt. This financial flexibility, coupled with a strategic partnership with Nvidia, positions the company to overcome execution and financing challenges, scale its GPU capacity and AI service revenue, and achieve significant ARR growth beyond 2025.

Catalysts over the next 12 months – including equity/debt financing for GPU capacity expansion, potential anchor customers, strategic asset sales or investments, and an IPO of its closest peer CoreWeave – are likely to drive a meaningful re-rating of the stock.

As of January 30, Nebius is trading at ~\$32/share. With a blended valuation of **\$130/share**, we believe there is **~305% upside**. Should the company scale similarly to CoreWeave and hit the CEO's target of 1GW in 2026, we believe a Bull Case of **\$249.50/share**, or **~685% upside** is possible. Additionally, the potential increase in per-share value from the company's plans for Avride/ClickHouse ranges from \$13/share to \$73/share exceeding Nebius's current share price by several multiples.



Catalysts:

- Accretive financing deals
- GPU capacity expansions
- Anchor AI customer
- DeepSeek clarification to NBIS's business plans
- Sell-side initiations
- Avride and non-core asset sales or investments
- CoreWeave IPO

I do not hold a position with the issuer such as employment, directorship, or consultancy.

I and/or others I advise hold a material investment in the issuer's securities.

Catalyst

- Accretive financing deals
- GPU capacity expansions
- Anchor AI customer
- DeepSeek clarification to NBIS's business plans
- Sell-side initiations
- Avride and non-core asset sales or investments
- CoreWeave IPO

Messages

Subject	Barriers to Entry?
Entry	02/06/2025 10:26 AM
Member	Weighing Machine Capital

Thanks AAOI for an interesting idea.

Curious if you could walk through the barriers to entry for what NBIS is doing.

How is NBIS not analogous to EQIX (with the difference being the compute/infrastructure inside the 4 walls is owned by NBIS)

What "edge" does NBIS have in the design of their architecture? Is there a sustainable edge with that?

Seems like NBIS could get hurt as the pace of GPU upgrade cycle keeps quickening. Doesn't that create an obsolescence risk with each upgrade cycle for what they've already built?

Subject	Re: Barriers to Entry?
Entry	02/06/2025 03:10 PM
Member	Shoe

That's my concern also. I've been trying to dig through Coreweave (CoreScientific) and Nebius too.

I think these are shorts actually. This is going to be an extremely commoditized business. It's the same as equipment rental companies, but far worse given how quickly GPUs depreciate, nvidia is on a 1 year product cycle, and how easy it is to shop for lower rental rates.

The rentals rates on GPUs are extremely commoditized and competitive. It's just a race to the bottom. It'll be like bitcoin mining

Unless you can build services and software on top of it, there shouldn't be much in excess profits, and this biz will be low margin and low multiple.

<https://cloud-gpus.com/>

<https://www.latent.space/p/gpu-bubble>

EQIX is better since they have co-location / interconnect revs, and the value is in the real estate. NBIS and Coreweave/CORZ are worse.

Subject	Re: Re: Barriers to Entry?
Entry	02/06/2025 03:25 PM
Member	gtwt

What keeps me from subscribing to this view (although I fully share the logic) is that the team running Nebius is actually top-notch. There seem to be so many easier shorts. Just my 2 cents.

Subject	Re: Re: Barriers to Entry?
Entry	02/06/2025 03:41 PM
Member	ahnuld

Shoe, I am like 99% certain core scientific is not coreweave. Coreweave did a series C last spring and hasnt sold any shares publicly

<https://www.prnewswire.com/news-releases/coreweave-secures-1-1-billion-in-series-c-funding-to-drive-the-next-generation-of-cloud-computing-for-the-future-of-ai-302133328.html>

Subject	Re: Barriers to Entry?
Entry	02/06/2025 03:47 PM
Member	AAOI

My pleasure Weighing Machine and happy to walk through the barriers to entry.

To start, I'd point out that the barriers here are knowledge of data center buildouts and capital. What our thesis shows is that there are few who can build a datacenter to perform as efficiently as the hyperscalers (see graphic in the write up). From our research, this PUE efficiency is not simply a function of cooler air in Finland. The company has not only a specialized way of clustering compute that makes it more efficient, but they have proprietary tech for co-generation that allows them to take excess heat and use that energy to spin up a turbine to transfer that power back into the DC, although some of the excess heat goes to homes nearby as they publicize. So, they're very efficient data centers with an advantage that is replicable and therefore can be pressed and leveraged over time.

Keep in mind, NBIS PUE for their HPC DC is right around GCP and Azure, which is telling. Ergo, with such an energy efficient low cost operating model they should be able to prioritize cost effective solutions that appeal to strongly budget conscious customers seeking the most value for their money. So even if I grant you that the next gen data center industry does ultimately get commoditized, NBIS's low cost moat/superior value proposition would still have the secret ingredients for outsized returns embedded within its differentiated operating model, to say nothing of its difficult to replicate longstanding relationship with the industries most important supplier and the preferential pricing that flows from it. There is also the company's geographic and regulatory advantages in Europe as well as its stable of best in class engineers, which isn't exactly easy to replicate from scratch either, but I'm sure you get the point.

Finally, we need to also consider that it takes a tremendous amount of capital to stand this stuff up, so access to markets is key and this is part of the reason NBIS wanted NVDA in on the latest round. At the moment, it helps on the cost of capital side with such confident backers.

With respect to your question on EQIX, do you consider GCP or Azure analogous to EQIX? NBIS isn't just building these DC's for low value and renting them out. They're building a cloud company from the bottom up specifically for AI workloads from AI customers. Note NBIS can do everything in house. They can build them, run them and above all, are in the process of putting a software service around it. Capacity is such that they're doing co-location at the moment but will be greenfielding new DCs in the US soon.

So as detailed in the thesis, NBIS is a play in two parts. The first is scaling as quickly as possible the raw compute necessary to be a leading AI cloud provider. The second is to wrap an open source software platform around that compute. By providing every possible company with ability to develop, run, and monetize AI applications, they are uniquely positioned to capture as much of the inferencing market as they can, the next leg here after all the training done to date.

This second phase is in a very early stage, its called AI studio, graphic presented in the thesis. Here's what we wrote *"In this second phase, Nebius aims to establish an open-source software ecosystem centered around its large-scale GPU infrastructure, offering cost-effective, full-stack AI solutions. **This strategy would not only more effectively monetize older node GPUs but also protect the company against competition based solely on access to the best GPUs and their supply and demand.** By integrating tailored AI solutions with hyperscale-level compute capabilities, Nebius can differentiate itself from hyperscalers, which often lack customized offerings. This approach also positions Nebius to outcompete the fragmented market of point solutions and secure long-term leadership in AI infrastructure and software."*

And further down *"With infrastructure and then services and software (either built internally by Nebius or through an API), a potential AI user would have a one-stop shop to build/use an AI, with pricing based on tokens. An "AI token" refers to the smallest unit of data that an artificial intelligence model processes, particularly in natural language processing (NLP), where text is broken down into individual pieces (such as words or parts of words) for the AI to understand and base its responses on. The tokenization model simplifies billing, allowing customers to pay only for the exact resources they use. Pricing is based on how many units of words or numbers one wants to analyze (e.g. 1,000 tokens per \$1). A full-stack AI service would simply price its tokens above the cost to run a token, playing into Nebius's low-cost capabilities."*

Tokens also provide greater flexibility and scalability, enabling customers to allocate resources as needed. For example, customers could increase their token usage during training phases and scale it back during maintenance. This would especially benefit startups and research institutions with dynamic AI workloads, driving significant cost savings by avoiding over-provisioning."

Integrating tokens into an expanded ecosystem – including elements such as AI tools, data labeling services, and software products – enhances its value proposition and deepens customer engagement, creating a wider moat than Nebius would have from infrastructure alone. Importantly, the economics of this ecosystem would disassociate AI customers from the underlying compute hardware, meaning customers would pay per unit of what they would like to compute versus paying specifically for the compute hardware, letting Nebius distance itself from the GPU-only infrastructure treadmill model."

Before I wrap, let's just assume all of the above is nonsense and the next gen data centre biz is a zero. That's silly but let's run with it regardless given it highlights the fact that NBIS possesses multiple ways to win thanks to its ownership of two unrelated non-core assets that together are likely worth ~2x NBIS's current price. Avride could realistically get sold to a strategic for the implied value of the enterprise in its entirety - and click house provides a nice cherry on top for all the reasons laid out in the piece. So even with draconian and completely irrational premises/assumptions, I think the probability of permanent capital loss at or around the recent price is extraordinarily low. For that reason we love the set up and think it's a pretty rad way to get exposure to a major secular megatrend without all the hype and risks associated with other ways to play it. Happy to add more color if I missed anything when I have some more time, but that's how we are thinking about it in a nutshell.

Best,

AAOI

Subject
Entry

Re: Re: Re: Barriers to Entry?
02/06/2025 04:39 PM

Member Shoe
I know, I'm just saying Corescientific is basically a bet on Coreweave and the discount rate you use on those cash flows.

Subject Re: Re: Barriers to Entry?
Entry 02/06/2025 04:40 PM
Member tahoe88

The point about being a more efficient operator of their GPU clusters I don't think tells the whole story. The business model works better if they can have lower input costs (i.e. lower PUE means lower power costs) but that still doesn't necessarily make it a good business model. I think this business model resembles something like bitcoin mining or Wework where you have a very capital intensive liability side mismatch with the revenue opportunity, which works great when things are good but turns quickly if demand gets competed away. I could be wrong, but I also think the hyperscalers are in a much better position to disintermediate these businesses over time if it turns out to be a good business model.

The PUE point is interesting, but again isn't the whole point. Bitcoin miners operate at 1.0x PUE's. Hyperscalers could probably operate at lower PUEs if they didn't have renewable energy targets (which they are quickly eschewing). PUE's in the Nordics are among the lowest in the world because of lower temperatures - I know you mention this, but I don't think its a coincidence. I also don't think the heat reusability concept is novel - I have heard about this with other Nordic DC operators.

For NBIS to win they can't have one DC cluster, they need to expand into numerous markets and I think they will get exposed. It's also unrealistic that they build their own data centers everywhere, so their PUEs will be more at the mercy of the landlords which may have much lower PUEs. I think this is important to understand - PUEs won't entirely be in their control unless they build their own data centers which is extremely unlikely because of the cost and time to power, so they will be required to lease space just like Coreweave etc.

The best thing going for this company seems to be management. Maybe they will figure out how to win and outcompete, but it doesn't change my opinion on the business model and differentiation compared to peers.

Subject liquidation value
Entry 02/06/2025 04:45 PM
Member differentiatedfractal31415

Can you elaborate on the assumptions underlying the liquidation value estimate you provided - given the debates posted here, thx

Subject Re: Re: Barriers to Entry?
Entry 02/06/2025 05:30 PM
Member JohnnyFinance

Shoe - Good response by AAOI, but my simple mind has a more succinct one that may be helpful.

My simple take (note - I have an extremely small position in NBIS) - NBIS's value is in (i) the scarce and valuable talent they have assimilated, (ii) their relationship with NVIDIA to obtain the (currently) scarce and valuable hardware they need to operate their business, (iii) the sticky and valuable customer base that their data center assets possess / can attract / can retain, and **(iv) the potential further monetization of these customers in a high margin and capital efficient way through NBIS's (potentially) highly scaled / economical, open source software ecosystem via service and maintenance contracts (customization, maintenance, etc. for their customers).**

NBIS wouldn't have to compete on rental economics (they could just break-even on the hardware and make money from monetizing their software).

The latter point (**bold above** - monetization of their open source ecosystem) is the most important - If NBIS can become what I'll call "the Red Hat of AI" or even be perceived as having the potential to become "the Red Hat of AI" by the market, then the company will likely be worth substantially more than its current enterprise value of ~\$5.7bn. The company owns a few misc. assets that give them a few additional "shots on goal" as well.

In any event, it is not something I would want to short. Just my two cents.

Subject Re: liquidation value
Entry 02/06/2025 06:39 PM
Member AAOI

Sure thing differentiatedfractal31415! The simple liquidation math is below and for various reasons I think you'll find the walk trollishly conservative.

The liquidation value walk is as follows...

1. \$2.6 billion in cash (\$3 billion pro forma equity raise less \$400 million 4Q capex)
2. Finnish Data Center - \$2.1 billion, 65 MW at ~\$30K per MW.
3. Avride: ~\$1 billion (2x invested capital).
4. Other Software: \$1,600 million (\$1,100 million for ClickHouse, \$300 million for TripleTen, \$200 million for Toloka).

Finally, some additional data points to reinforce how sandbagged the above assumptions are:

Avride: With the exception of Baidu and Waymo, there isn't a comparable autonomous asset on earth that has achieved comparable milestones or as much autonomous mileage as Avride has to date, save Tesla. This is a business operating in high growth markets, a strategic partnership with Uber and as such, would be a great tuck in acquisition for Uber, TSLA and/or Waymo. So while I won't go to the trouble of repasting the chart and paragraphs from the thesis, rest assured if you're reflexively discounting the value of this asset because you've never heard the name Avride before you're a fool.

Mind you, we've valued it at \$1bn to be intentionally trollish in the liquidation assumption above, not because it's an accurate reflection of its economic value on an absolute or relative basis in light of its scale, marquee partnerships, ~50k passenger rides, ~200k robot deliveries to date, and so on. After all, with a 1.2 trillion market cap, I'll leave it to others to put a value on Tesla's autonomy related optionality, and we know Waymo was recently valued around \$45bn. Cruise and Zoox are a joke - and Baidu isn't really relevant here for obvious reasons. Nonetheless, one doesn't necessarily need a scale to know a man's fat. Likewise, if we value Avride on an equivalent autonomous mileage basis as Waymo, that implies a value approximating ~30% of Waymo or ~\$14 billion vs. an NBIS EV of ~\$5bn for the whole enchilada. Let's apply a discount to that estimate above and run with 20% of Waymo's value or \$9bn just to be safe. See what I mean?

Yeah, \$1bn probably isn't the right figure here but members can roll up their sleeves and do the work for themselves.

Clickhouse: See the press release linked below. You may not have heard of it before but rest assured that's the feature with this idea, not the bug. Google it on X. Google it on Google. This isn't some shitco stub worthy of a garbage barge, that much I can say for sure.

<https://clickhouse.com/blog/clickhouse-announces-strategic-collaboration-agreement-with-aws-to-advance-real-time-data-analytics-and-generative-ai-innovation>

Hopefully that helps get you started.

AAOI

Subject	Lets do better, can we?
Entry	02/06/2025 06:48 PM
Member	AAOI

So guys, I'm picking up on a trend here. It's almost like no one has really read the thesis as I'm basically already way ahead of you on questions and I'm just copying and pasting from the doc. You can see this via the degradation in the avg quality of the tags in recent years, as well as the multitude of 1's and 2's certain members have been pinging this idea since I posted it, yet until today not a single intelligent question or counterargument was posed. As far as the tags, I mean at least they used to be clever and funny when - perhaps especially when - they were aimed directly at me. Now it's just kind of sad as I'd say that 95% of the time they are equal parts embarrassing and dumb.

SToCK prOmOTIon!?! WhY AlWAYs PuMpInG!?! Are you serious? Lmao! Is this really the best we can do?

At least the marcellus tag makes a decent attempt to be clever, but even there it's clear the author didn't actually read a word of what was written here, otherwise I doubt it would have been done. Or perhaps their reading comprehension is just bad. Who know's, but let's have some fricken self respect when lobbing a proper bomb. Maybe strive for something that elicits something a bit more than a bewildered "huh?" if you're going to put forth the effort. Make me laugh...be uncharacteristically mean...something...anything really but these tired ass and frankly embarrassing tags. Which is a long winded way of stating, let's do what we can to make VIC tags great again as these wack, completely nonsensical jabs ain't it. That much I assure you.

And while I'm at this rant, for those of you that gave me a 1 or 2 rating at least do me the favor of lighting me up along the way in a manner that explains why you think I'm wrong - if only out of respect (see my commentary on the ASTS short thesis, for example). I mean let's debate! This isn't personal.

Better yet, imo it's actually way better when someone takes the chainsaw out, both because if I'm wrong I want to know it - and because happy warrior thrust and parry exchanges like these (where people think I'm somewhere between a moron and dead wrong) makes the whole exchange a lot more fun. This is fricken VIC for God's sake.

Subject	Re: Re: Re: Barriers to Entry?
Entry	02/06/2025 06:51 PM
Member	AAOI

Beautifully stated JohnnyFinance. Agree with every word.

Fwiw, I'm more than happy to just let the bears go wild with confirmation bias and back slaps after a whole week of reading DeepSeek articles. Lmao.

Subject	Re: Re: Re: Barriers to Entry?
Entry	02/06/2025 06:53 PM
Member	AAOI

Great stuff Tahoe, let me discern a worthy response and circle back tomorrow when I have more time to offer some thoughts and counterpoints.

Subject	Re: Lets do better, can we?
Entry	02/06/2025 06:54 PM
Member	Weighing Machine Capital

AAOI,

Is the CEO locked-up on any of his shares he owns?

What was the purpose of the Finnish DC when it was greenlighted and in what year was that?

Thanks

Subject	Re: Re: Re: Re: Barriers to Entry?
Entry	02/06/2025 07:14 PM
Member	JohnnyFinance

Haha - DeepSeek, or as I like to call it, DeepHideAndSeek!

Subject	Re: Re: liquidation value
Entry	02/06/2025 08:09 PM
Member	Ray Palmer

Can you help me with the \$2.1B valuation for the finnish data center? \$2.1B on 65 MW = \$30 million per megawatt? That seems really high to me. I'm doing this from memory as I was much closer to the industry over the fall, but I think brand new AI data centers with hundreds of megs of power (and more megs are important; there are returns and benefits to scale.... 65 MW is honestly kind of small) would cost like \$10M/MW.

Subject	Re: Re: Re: liquidation value
Entry	02/06/2025 08:14 PM
Member	Shoe

Depends on the networking and complexity of the data center.

These days people cite \$15-20mm / MW.

\$30mm / MW would have to have to much snazzier equipment

Subject	Re: Re: Re: Re: liquidation value
Entry	02/06/2025 10:45 PM
Member	tahoe88

Build cost probably \$10-12mm/MW + GPUs / equipment usually 3x the build cost. Seems reasonable if including equipment, too high if not.

Subject	Re: Re: Re: Re: Re: liquidation value
Entry	02/07/2025 11:22 AM
Member	AAOI

Tahoe88 is correct, it's including equipment. We included the \$400m purchase of new GPUs in 4q as they're going to the Finland data center.

Subject	Re: Re: Lets do better, can we?
Entry	02/07/2025 11:28 AM
Member	AAOI

Not sure on his shares being locked up but he certainly isn't selling his stake anytime soon. I'll check and report back but simplistically, if NBIS wasn't the second chapter in his life's work I doubt NBIS would exist in the first place; he'd be on a beach somewhere fishing. Arkozy is definitely the right type of crazy.

The Finland DC was originally built in 2015 to handle Yandex's search engine and cloud services. Hope that helps!

AAOI

Subject	Re: Lets do better, can we?
Entry	02/07/2025 11:55 AM
Member	yarak775

I agree that much can be done better on VIC...but there's some serious irony in getting on your high horse re: NBIS when the idea was just posted here 2.5 months ago...

Subject Re: lockup question
Entry 02/07/2025 11:57 AM
Member yarak775

He's not locked up - nor should he be as this was not an IPO. Agree w/ AAOI tho, he's pot committed here. He could have just sold the whole thing and walked away with a billion dollars plus.

Subject Re: Re: Re: Barriers to Entry?
Entry 02/07/2025 12:14 PM
Member AAOI

Tahoe, on your DC comments. We're well aware of Nordic dynamics. We're sold a bunch of puts on APLD near the lows for similar reasons. Colder temps are better. Now I think we're likely ahead of you on the cooling tech. Broadly speaking, it sounds the same as other DCs that do energy regen, but after talking with someone familiar w their tech in extensive detail, they do think it's a special sauce that makes it differentiated/replicable. And again, with where the stock is we think we can be wrong on a lot of it and still come out way ahead. Your points are valid though. Great comments.

I'd also just ask have you guys done any digging on what coreweave is earning right now as far and margins etc on this terrible GPU rental business? We think gross margins should be above 80%. And some easy inference puts EBITDA margins in the 50-60s.

Clearly the market likes AI capex. And NBIS is levered to that. Better yet, if the massive increases in AI capex guidance this week in the wake of the Deephideandseek freakout isn't telling, not sure what could be. In any case, the market is such that they're likely going to be upping guidance in a big way. Plus, if Coreweave IPOs with numbers anywhere close to what we laid out above it will be an easy comp for people. The upward reverse gravitational pull from these dynamics should be substantial, so even if time proves your concerns correct, do you really want to be in front of all that with an accurate analysis of where the GPU market may be in 3 years? Additionally, as you allude to, is this really the type of management team you want to bet against? Arkozy is an absolute beast.

Point being, there are so many non correlated ways to win at the present valuation the risks you point out are ones I'm willing to take. Thoughts?

Subject Re: Re: Lets do better, can we?
Entry 02/07/2025 12:23 PM
Member AAOI

Yarak, totally hear you on that, which is to say there was a reason I chose specifically not to critique the "not a new idea" tag, which was fair. Thank you for bringing attention to the name as we've always been big fans of your work (Hilariously enough, your first NBIS post beat me by literally hours so I pulled it the first time around). At any rate, after reading the comments on the original thread regarding deepseek it seemed like no one was going to discuss the idea as a long, so we dropped an updated and, in our view at least, a particularly timely thesis specifically to refute the deepseek claims (given the 40% cliff drop last week).

AAOI

Subject Re: Re: Re: Re: Barriers to Entry?
Entry 02/07/2025 12:32 PM
Member Shoe

Yes gross margins and EBITDA are high , but D&A costs are also high very high and you can't ignore them.

I'd expect their financial model to look like equipment rentals (e.g. URI) [or aircraft lessors] with similar financial metrics. These don't get high multiples

Subject Re: Re: Re: Re: Barriers to Entry?
Entry 02/07/2025 01:53 PM
Member tahoe88

Fair enough and not necessarily saying its a short. Just saying I still have yet to see someone in detail outline the unit economics and ROIC over a sustained period of time and my intuition is that its probably good right now because we're in the early period of the AI adoption curve, but expect that to get competed away over time, likely by the hyperscalers who could recreate this business very easily. Coreweave is selling a commodity - compute - which like other telecom asset in history is typically deflationary. Thats my sense as a non-technologist. Coreweave I think has its own challenges particularly one large contract by a hyperscaler that will dictate its destiny - my concerns apply to them equally.

I also think the hyperscaler capex guides are being overlooked. What they are talking about in their 2025 guides are to capitalize leased and built assets that were 2023 decisions. MSFT for example is spending a lot more on GPUs than building new data centers, which again is a lagging indicator. The leading indicator is data center leasing / land purchases because of lead times, and again MSFT pulled back from something like 2 GW of commitments over the last two quarters. Of course the others are picking up the slack (notably Oracle on behalf of OpenAI), but I don't the hyperscaler capex situation is as universally bullish as it seemed.

Subject Re: Re: Re: Re: Re: Re: liquidation value
Entry 02/07/2025 02:11 PM
Member Mason

I think a great write at AAOI. On the liquidation value, so we are assuming for Finland, \$10m per MW of infra costs so \$650m, then the \$400m of GPU costs they spending in 4Q so \$1.05bn, so we have \$1bn remaining from your \$2.1bn? They have \$1bn of GPUs already at the site? Thanks.

Subject
Entry
Member

Re: Re: Re: Re: Barriers to Entry?
02/07/2025 02:15 PM
Weighing Machine Capital

Everything is about the details, right?

Yes, Datacenters are being built left and right and center. But, need to distinguish between the Gigawatt-sized datacenters that the hyperscalers have in the works. These are being built to house 100s of thousands (if not million+ GPUs) in order to PRE-train and POST-train frontier models (in the race to AGI).

Datacenters like the one that NBIS owns is obviously not that and was built a decade ago which means there's some outdated infrastructure that needs to be replaced to bring it to the cutting edge of datacenters being greenlighted today.

And this one is capable of inference, not training models. There's literally 1000s of datacenters the scale of NBIS' Finland datacenter that are being constructed right now SPECIFICALLY for inference. Heck, Meta exclusively uses shittier AMD MI300X GPUs for INFERENCE. And as we saw with Deepseek, the biggest takeaway was that there is substantially less barriers to enter for INFERENCE, but not, FRONTIER model training.

I don't understand the special sauce that NBIS has? As a long term Yandex investor, I would say that Arkozy is good, but not worth the hype. Yandex won in Russia more because of the government's disdain for Google than anything else.

Subject
Entry
Member

Re: Re: Re: Re: Re: Re: Re: Re: liquidation value
02/07/2025 02:30 PM
Mason

Looking at the liq analysis again, looks like you're using the more positive scenarios you outlined for the non-core assets. Have you guys been able to get any more comfortable with those valuations outside of what you put. Using 2x 2021 valuation for a software biz a bit aggressive and then 2x invested in the ADAS biz I can understand but there is a big graveyard of companies in that space as well. Just trying to get more comfortable with the downside. Thank you.

Subject
Entry
Member

Re: Re: Re: Re: Re: Re: Re: Re: liquidation value
02/07/2025 03:06 PM
AAOI

Mason, if you ask around you can find out about clickhouse. Apparently it's been the best performing asset out of the involved VCs 2021 vintage assets. Thinking here is the company is getting ready for a raise following the AWS contract. \$4B is what some are thinking at the low end with that relationship in hand. Another way to approach triangulating this comes straight from the horses mouth. In interviews Arkady has explicitly stated "if I want to fund our guided GPUs purchase for this year I can just sell our clickhouse stake," so $1-1.5b / .28 = \$4.4b$. There are other reasons we think this is conservative but we feel pretty good that said mark would be realized if they were to put the asset up for sale today.

AAOI

Subject
Entry
Member

GPU Economics
02/10/2025 02:36 PM
RugPull1995

Can you explain how NBIS GPU economics are ROIC positive? I have a basic model for IRR on renting out NVDA GPUs from conversations with Llamda Labs and Coreweave on total capex per GPU, opex / EBITDA margins (was 70-80% when H100 pricing was \$4 but now certainly less), and then assuming price degradation on Blackwell that we saw on A100 and H100 (A100 went from \$2.50 to now low 1s in two years, H100 went from \$4 to now \$2 / hr in two years, so roughly 50% cut in two years seems to be deflation in GPU pricing), and the math doesn't pencil out. The only caveat where the math could pencil is if you manage to lock a lot of customers into LT contracts on initial pricing of Blackwell GPUs, but Nebius has themselves admitted that they don't currently have that strategy of LT contracts, unlike CW and Llamda Labs.

I guess explain where my assumptions below are wrong?

Cost / GPU	42,000	<3M ASP for NVL72 (\$3M / 72 GPUs = \$42K per GPU)							
Other Capex / GPU	20,000								
Total Capex / GPU	62,000								
GPU / Hr	\$ 3.50								
Hrs/Yr	8,760								
Utilization	80%								
Hrs Util.	7,008								
GPU Useful Life	8 Years								
		1	2	3	4	5	6	7	8
	Y0	Y1	Y2	Y3	Y4	Y5	Y6	Y7	Y8
Hours	7,008	7,008	7,008	7,008	7,008	7,008	7,008	7,008	7,008
(x) \$ Per GPU Per Hr	\$ 3.50	\$ 2.63	\$ 1.97	\$ 1.77	\$ 1.77	\$ 1.77	\$ 1.77	\$ 1.77	\$ 1.77
Revenue Per GPU	24,528	18,396	13,797	12,417	12,417	12,417	12,417	12,417	12,417
Memo: Cum. Revenue / Capex		0.4x	0.7x	0.9x	1.1x	1.3x	1.5x	1.7x	1.9x
Operating Expenses / GPU		6,570	6,570	6,570	6,570	6,570	6,570	6,570	6,570
EBITDA Margin		73%	64%	52%	47%	47%	47%	47%	47%
EBITDA / GPU		17,958	11,826	7,227	5,847	5,847	5,847	5,847	5,847
Cash Flow	(62,000)	17,958	11,826	7,227	5,847	5,847	5,847	5,847	5,847
MOIC		0.3x	0.5x	0.6x	0.7x	0.8x	0.9x	1.0x	1.1x
IRR		1.9%							

<Solved for from info given by 2 Llamda executives who said 75% EBITDA margins at 75-80% utilization on GPUs

Subject Re: GPU Economics

Entry 02/10/2025 02:36 PM

Member RugPull1995

Sorry - other thing I left out is conversations with multiple customers of NBIS who are saying initial Blackwell pricing / hr is being quoted at \$3.50 / hr

Subject Re: Re: GPU Economics

Entry 02/10/2025 04:49 PM

Member Weighing Machine Capital

This is looking more and more like a good short.

- No barriers to enter whatsoever
- NBIS does not have any scale and DC is already antiquated (Finland DC built 10 yrs ago)
- What software (SAAS or IAAS) does NBIS pretend to have/be building?

Subject Re: Re: Re: GPU Economics

Entry 02/10/2025 05:07 PM

Member Shoe

Also, these have like 4-5 useful lives if air cooled.

Supposedly liquid cooling extends that a few years. And submersion cooling even longer

Subject Only Bait In Town

Entry 02/18/2025 10:33 AM

Member yarak775

I wrote this last week, hesitated to post it, but seems more germane than ever, so here goes...

NBIS fits an investing framework I call "Fins." This might belong more in the topics section of VIC than in the messages, but NBIS is such a stunning example of Fins that I feel compelled to share it here.

Throw out your copy of Graham & Dodd, pour a cold beer, and bear with me, VIC friends, as I explain a concept based on the writings of Buffet (but not Warren.)

The late Jimmy Buffett wrote and recorded a reasonably well-known song called "Fins" in 1979. It's about a woman seeking peace and quiet in a warmer climate, but due to the large quantity of males and the paucity of females, she finds herself besieged by attention from would be suitors.

"Can't you feel them circling honey? Can't you feel them swimming around? You've got fins to the left, fins to the right, and ***you're the only bait in town.***"

A Fins situation in the stock market is similar. It starts with a theme or industry niche that is hot, but with very few ways to play that theme directly in the stock market. So, the single name (or few names) that are directly tied to said theme get a great deal of attention. ***They are the only bait in town.***

Historical examples off the top of my head:

TDOC - Telemedicine

TRUP - Pet insurance

TWOU - Online higher education

FLUT/DKNG - Sports betting

TSLA - Electric cars

ACLI - Inland barge shipping

FOXF - Premium auto/bike components

FRPT - Pet food

LEU (USEC) - Uranium enrichment

I could list dozens, but you get the idea.

1. Some commonalities of Fins situations: Current fundamentals do not matter much. Fins are generally emerging business models, concept companies, or established companies hoping to capture a super cycle. People aren't buying the steak; they are buying the sizzle. Or maybe even the aroma from the kitchen. It's all on the come.
2. There will be some metric that matters like orders or revenues. But all the big results - especially the profits - are usually several years out. It can neither be proven nor disproved conclusively that they will ever arrive. They get bid to very high valuations due to enthusiasm for the trend and scarcity value of the stock. They become popular, momentum, IBD50 type names.
3. Fins stocks also become popular shorts because investors (often too smart for their own good) question the actual quality, durability, or profitability of businesses. These questions, coupled with premium valuations (see #1 above) makes betting against Fins especially irresistible for dyed in the wool short sellers (I cast my lot with these poor souls, and have the scars to establish my bona fides.)
4. Eventually, the battleground is resolved. The weighing machine takes over from the voting machine as fundamentals present themselves. **This often takes longer than we would expect.**

TDOC is a great example of all the above. We identified TDOC as a low-quality player in a low-moat and increasingly crowded industry shortly after its IPO in 2015. The stock went parabolic during COVID, but even barring that, it remained frustratingly strong in the face of suspect fundamentals and high short interest for years. It was the only bait in town for telemedicine.



NBIS fits the bill for a Fins stock in spades. It is one of the only pure plays in the red-hot subsector du jour (AI data centers.) One of the largest comparable companies, CoreWeave, has seen its valuation skyrocket in private markets. Myriad other factors such as the former YNDX CEO, the special NVDA relationship, the Waymo-esque division help support a SOTP valuation help fan the flames of excitement.

On the other side of the ledger, we really have no idea what the economics of this business will look like in 2 to 5 years. As RugPull models, this may be an enterprise that simply runs in place, constantly having to pay up for the latest NVDA chips to outrun a steadily and steeply eroding pricing curve for older offerings. This can be postulated with math today, but it can't really be proven yet. The steak may be terrible when it arrives at the table but the sounds and smells from the kitchen still entice the patrons to order an entrée.

My guess is, like many Fins situations, the stock will go higher before it goes lower. Management has every incentive to get the stock price higher, they will be among the first in line for Blackwells, and the bright future the bulls will paint for this business can't be disproved in the near-term. I would not be inclined to short it at this time. Absent more DeepSeek type events, I think the bias is to the upside.

In 2 or 3 years...who knows? We will see several more cards over that period that will help us understand the profitability and durability of the business. We will likely get a CoreWeave IPO, which I do not think is necessarily a positive. Perhaps there is a valuation discrepancy that highlights an attractive opportunity in NBIS shares. But given CoreWeave's relative scale and profitability, I think this will take flows away from NBIS. Remember, *NBIS only gets attention because she is the only girl in the bar*. The sharks who swim on land will quickly turn their attention elsewhere if another (and especially a prettier) girl shows up.

NBIS: the only bait in town. Likely to go higher for the next 2-3 years, then likely a short from (much?) higher levels as reality sets in on the economics for a non-hyperscaler provider of commoditized services.

Post Script: Despite its popularity at parties and concerts, Fins is actually a pretty cheesy and terrible song. If you want to get into listening to Jimmy Buffett (which I strongly recommend), please don't start with Fins (or any song you've ever heard played over the speakers at a stadium or restaurant.) Start with: A Pirate Looks at Forty, Death of an Unpopular Poet, Tin Cup Chalice, Come Monday, Manana, Havana Daydreamin', He Went to Paris, Trying to Reason with Hurricane Season, Little Miss Magic, California Promises, and Migration to get an appreciation for Buffett's songwriting and overall vibe.

Subject

Re: GPU Economics

Entry

02/18/2025 10:42 AM

Member

gtwt

I think confirming or disproving RugPull1995's IRR math (post #31, basically suggesting zero ROIC for the NBIS business) would be incredibly helpful here. I haven't done the math myself but the numbers and logic make sense. Does the idea author or anyone else have counterarguments?

Subject

Re: Re: GPU Economics

Entry

02/18/2025 11:10 AM

Member

agentcooper2120

I'd chime with two points. 1) have heard from a few people the NBIS NVDA special relationship amounts to Blackwell GPUs at \$25-\$30k each (one source said as low as \$20k)

2) the lower rates such as the \$3.50 are for the first 1k hours and then bump up 25% afterwards.

Also to Shoes point, read a white paper that showed useful life of a GPU at full utilization was 80% longer when liquid cooled vs air cooled.

Subject

Re: Only Bait In Town

Entry

02/18/2025 01:31 PM

Member

RugPull1995

You might be correct, and personally I wouldn't short this for the time being given we appear to be in a period where retail is sending these stocks higher and higher, and I do appreciate you being intellectually honest that going long here is basically taking a punt, but think also important to distinguish between that and an actually good **investment** thesis. I've completed a bunch of work here and pretty much disagree with every fundamental assertion the author made - there is no software moat here and no differentiation - just talk to any of their customers and they will tell you price and availability is all that matters to them and they all use multiple neo-clouds to instill pricing discipline. Customers will also tell you that inferencing will not occur on the neoclouds - models will be trained on them but the use of the models at the application layer will be in SaaS companies / etc. that have their infra in the cloud (just ask the customers, their expenses are in the neoclouds and their revs from API calls from the hyperscalers)

It may take a long time, but as soon as there is no GPU shortage or VC funding dries up for these AI model builders that have no chance of success, these neo-clouds will be toast. To make money on this, you will have to see the peak coming before everyone else and get out, which to me is a hard game, but could work if you trade the stock well.

Just my two cents,

RugPull1995

Subject	Re: Re: Re: GPU Economics
Entry	02/20/2025 09:37 AM
Member	Shoe

They said it on the call. The GPU rental business will be commoditized. This is a bet on being able to sell other software and service on top of the bare metal in order to make an ROIC / margin. And they don't have any obvious full stack services yet.

:

"In turn, this gives us a better and enhanced way to create a return on capital invested and also creates new revenue streams for the company that enable us to command software-like margins. **So we're very excited about, from the top to the bottom of the stack, all of the innovation that Nebius is bringing to market. Going forward, we think pure bare metal offerings will become more commoditized. And it's our investments now across the full stack, especially at the top end of the stack, that are going to pay the biggest dividends in the mid- to long-term.**"

"in addition, we are a full-stack business, and as a full-stack business, we have additional value-added services that we can offer that could really help accelerate the ROIC as that software layer is a high-margin element to our business. It's still very small today, that software piece, but we think that as we grow our customer count and we start building more tools and software and services, we think that can actually be a nice contributor to revenue and margin and help with that ROIC"

Subject	Re: Re: Re: Re: GPU Economics
Entry	02/20/2025 10:46 AM
Member	Weighing Machine Capital

@Shoe,

Has NBIS disclosed anything tangible regarding this "full-stack" software?

In the new world of AI, where software is being mostly commoditized; companies with **unique data** that build **non-commoditized** software are the only ones going to command high margins.

The rest will be run over by the flywheel of better data --> better software --> more users --> more data --> better software --> more users...

NBIS "claims" their full-stack software is great. I beg to differ. They don't have any customers yet to provide them the data to make that software great!

At their sub-scale, it's going to be very difficult to chase the horses down that are way ahead of them in my opinion.

Subject	Re: Re: Re: Re: Re: GPU Economics
Entry	02/20/2025 10:53 AM
Member	Shoe

There is some software, but nothing meaningful significantly value add above helping to manage the core bare metal.

This is a VC bet that they'll be able to move up the value chain, which I don't see

Subject	Re: Re: Only Bait In Town
Entry	02/20/2025 02:33 PM
Member	yarak775

RugPull - I largely agree with your two cents. And I'm not really advocating owning NBIS at this point...just trying to share some pattern recognition based on 20 years of generally shorting Fins stocks too early.

The way I view it, by their nature Fins are no longer investments. They are...pick your descriptor...battle grounds, trading vehicles, sardines...generally a high tension tug of war between two forces of the investing world. While I do think there are ways to trade and profit on Fins, this is Value Investors Club not Value Traders Club, so I'll spare everyone another musical metaphor.

Subject	Re: GPU Economics
----------------	-------------------

Entry
Member

02/24/2025 06:42 AM
ahnuld

just figured id share, this is a slide from a seed stage AI infra company I am diligencing. Slide is from August.

Market is saturated on the hardware layer > \$5 Billion in competing investments

GPU Model	Rental / hr (assume 100% uptime)	Electricity Cost (Assuming 0.16/kwh)	8xGPU node cost (include setup, network & 3 year labour)	ROI (assuming no price drop)	Normalized ROI (assuming price drop by 50% across lifetime)
H100 SXM	\$2.85	\$0.16	\$400,000	2.10 years	3.15 years
H100 SXM	\$1.00 (limited supply)	\$0.16	\$400,000	6.23 years	9.89 years
A5000	\$0.28	\$0.05	\$40,000	2.42 years	3.63 years
L4	\$0.20	\$0.02	\$35,000	2.76 years	4.14 years

Unless you have nearly free electricity
Or can charge a higher rental rate with a Sovereign AI play

Buying new hardware, and running it at market rate,
has little to no ROI in < 4 years

Subject Coreweave S-1
Entry 03/03/2025 05:27 PM
Member RugPull1995
Anyone who is reading this S-1 right now and thinks this is positive for NBIS... you live on a different planet than I do.

Yours,

RugPull1985

Subject Re: Coreweave S-1
Entry 03/03/2025 08:45 PM
Member Ray Palmer
Could you elaborate?

On the one hand, I don't disagree with you- Coreweave is obviously way more scaled, sophisticated, etc.

On the other hand- doesn't the coreweave S-1 show that if NBIS is successful, there could be a pot of gold at the end of this rainbow?

Subject Re: Coreweave S-1
Entry 03/03/2025 10:13 PM
Member Shoe
I think the economics are OK and inline with the ballpark of what I thought it would be.

In 2024 they did 65% EBITDA margins, and 15% EBIT margins (which includes 1% of SBC)
In Q4, they did 60% EBITDA margins and 15% EBIT margins

So overall, not bad

On Q4 run-rate, this is 2.5% EBIT/assets, 11% EBITDA/assets.

I think this level of margins and returns makes sense for what is basically an equipment rental business. And they're still growing a lot, so there's room for margins to improve from here.

Subject Re: Re: Coreweave S-1
Entry 03/04/2025 10:22 AM
Member Shoe

However, their useful life may be overstated. Most hyperscalers are using ~4-5 years on technology depreciation

The estimated useful lives of the Company's property and equipment are as follows:

Technology equipment	6 years
Software	3-6 years
Data center equipment	8-12 years
Furniture, fixtures, and other assets	3-5 years
Leasehold improvements	Shorter of remaining lease term or estimated useful life

Subject Re: Re: Re: Coreweave S-1
Entry 03/04/2025 12:09 PM
Member RugPull1995

Shoe - to say their useful lives are overstated is an understatement. A few points on why this S-1 is not good - personally I think this IPO gets pulled, and if it does come to market trades very poorly.

- GPU useful lives are not 6yrs - can debate this point back and forth forever - just don't believe it. A100 and H100 pricing got cut in half in 2yrs, so at the very least they should be depreciating GPUs on an accelerated basis, not a linear basis given the revenue potential of the GPU is massively weighted to the first 2 years. Yes they have some long-term contracts with MSFT (more on that in a bit) on Hopper when the world had a huge GPU shortage, but how long can they do that for. The idea that FCF burn is \$6B and they are Adj. EBIT + is laughable in my view. Also datacenter equipment 8-12yrs? What? Networking equipment for example is well known to have 4-5yr cycle lives.
- MSFT is a 60%+ customer and 2nd customer is 15% so top 2 customers are 80%...ok... so your largest customer is Satya Nadella who literally went on a podcast last week and talked about how GPUs will get overbuilt and how happy he is to be a leaser and not a buyer since it is easier to let leases roll off
- MSFT can cancel the master services agreement "for cause" in 45 days - I am not sure what cause is defined by cause but incremental negative
- Founders sold \$500m of stock leading up to the IPO
- Just generally speaking - if the payback period on GPUs is 2.5-3yrs like they said in the S-1 and Nebius said on its call - then assuming GPU prices you can get in years 4-6 are the same as years 1-3, your MOIC over 6 years (theoretical useful life of GPUs) is 2x... forgetting about time value of money - why is this EV worth more than 2x book (I get they use leverage to amplify equity value, yes) - invested capital here is \$9B and last round was raised at \$30B

Subject ehem
Entry 03/04/2025 12:18 PM
Member mojouris
 Coreweave is WeWork...?!?

Subject Re: Re: Re: Re: Coreweave S-1
Entry 03/04/2025 02:09 PM
Member lars

Don't forget the material weakness in their controls. Do you know what the target timeframe is for the offering? I think it could be a pretty material negative catalyst for AI / growth stocks when this gets pulled.

Subject Re: Re: Re: Re: Coreweave S-1
Entry 03/04/2025 02:44 PM
Member lars

Also, NVIDIA is an investor in CoreWeave....they're an investor in Cohere, a CoreWeave customer (who's not profitable, as far as I know)... Outside of the LLMs that have other businesses, this has a very house-of-cards feeling - is it ok that NVIDIA is funding customers who are burning cash?

Subject Re: Re: Re: Re: Re: Coreweave S-1
Entry 03/04/2025 02:49 PM
Member gtwt

There are good insights into this: <https://x.com/kashyap286>, <https://x.com/kakashiii111>

Subject Re: Re: Re: Re: Coreweave S-1
Entry 03/04/2025 03:47 PM
Member Shoe

I hear what you're saying and agree to an extent, but I don't think this isn't nearly bad enough to be pulled. I was short WeWork bonds around the IPO, that was definitely bad enough to be pulled.

They're profitable, the FCF burn is from growing a lot.

Everyone knows MSFT is a big customer. MSFT is saying they'd rather lease, not buy, not to get out of the lease, but because they think prices go down.

But yea I think this is a bad business, it's someone between bitcoin mining and construction equipment rental

Subject Re: Re: Re: Re: Re: Coreweave S-1
Entry 03/06/2025 08:08 AM

Member

lars

Probably not the press the bankers were hoping for - <https://www.ft.com/content/f3d9d339-42ef-4979-bf52-89ecd699dea2?Fds-Load-Behavior=force-external>

Subject

Re: Re: Re: Re: Re: Coreweave S-1

Entry

03/06/2025 04:52 PM

Member

Shoe

Coreweave denied it but where there's smoke there's fire?

Who knows these days, the media has been blowing stuff out of proportion for clicks lately

Subject

Re: Re: Re: Re: Re: Coreweave S-1

Entry

03/06/2025 05:22 PM

Member

lars

I'll take the FT over folks trying to sell you an IPO with a 62% customer and a material weakness in their internal controls after they've just unloaded \$500m themselves.

Subject

Interesting write up on Coreweave

Entry

03/18/2025 11:11 AM

Member

RugPull1995

Think >50% chance IPO gets pulled. From what I am hearing they seem desperate to go public and insiders are desperate for liquidity, so possible they cut the valuation, downsize the offering, and shrink the float with insiders participating to jam it through.

Will be an amazing short opportunity if it manages to squeak through.

If even 25% of what is written in this article is true, this is an absolute trainwreck.

<https://www.wheresyoured.at/core-incompetency/>

Subject

Re: Interesting write up on Coreweave

Entry

03/18/2025 11:26 AM

Member

lars

NVDA invests in CRWV, CRWV takes that money and buys NVDA GPUs, and then NVDA rents those GPUs back from CRWV.....what could go wrong??

Subject

Re: Re: Interesting write up on Coreweave

Entry

03/20/2025 07:30 AM

Member

lars

Looks like they're already lowering their sights from \$4b to \$2.3-2.7b. Those proceeds plus their total existing liquidity wouldn't even get them to last year's capex number. I think this one is going to get ugly. If this doesn't get done, which I agree with RugPull1995 is <50%, it will be interesting to see how this gets addressed in the private markets.

<https://on.ft.com/4bFyS2Q>

Subject

Re: Re: Re: Interesting write up on Coreweave

Entry

03/20/2025 09:52 AM

Member

onodacapital

I think there's more demand than people think. Could be wrong ofc. They also signed a \$12B deal w/ OAI, usually get 20% upfront = \$2.5B. If they raise \$2.5B in IPO - which I think they can, the question is if they'll take the price - that's \$5B total. Think there could be more OAI deals coming too. EBIT should be >\$1B this year. I wouldn't own it at the prices they're talking, but I wouldn't write it off as a zero yet - esp if they get a halfway decent IPO.

Subject

Re: Re: Re: Re: Interesting write up on Coreweave

Entry

03/20/2025 10:01 AM

Member

lars

I would be concerned that you're being way too generous re: the OpenAI deal, for two reasons:

1. When a contract value is characterized as being "up to" a headline number, the lower bound is often zero / as low as the customer wants.
2. The deal was announced seemingly in reaction to the MSFT news and came with an equity stake for OpenAI - looks quite desperate to me.

Subject

Re: Re: Re: Re: Re: Interesting write up on Coreweave

Entry

03/20/2025 10:27 AM

Member

onodacapital

Anything could happen, but that's not true here. Lots of people involved speak to the legitimacy of the contract. "MSFT pulling back" news is also well known that OAI isn't pleased w MSFT speed / execution in DCs. I'm sure some truth in what Satya says, but also lots of posturing.

Subject

Re: Re: Re: Re: Re: Re: Interesting write up on Coreweave

Entry

03/20/2025 10:40 AM

Member

lars

Agree to disagree - I don't see how you can say "that's not true here". It could well be true, especially given the OpenAI deal timing & equity component as a signal. This company seems desperate to go public a week before more tariff fireworks go off, and the downsized raise they're talking about doesn't even get them off of the funding treadmill for 12mos. That in the context of them essentially buying from one supplier and selling to one customer of any materiality (maybe two, and ok even three if you're going to count their biggest supplier and also investor), along with the laundry list of other red flags, and I don't know how anyone with a fiduciary duty can rationalize putting capital at risk in this equity at any practically any price.

And as I've said before, the sketchy fact pattern is not just at CRWV, but across the ecosystem - NVDA invested in CRWV, CRWV took that money and bought NVDA GPUs, and then NVDA rented back a portion of those GPUs as CRWV's second largest customer. Furthermore, NVDA invested in multiple other cash-burning CRWV customers (who are also, at least indirectly and very likely directly, NVDA customers), including Cohere and Mistral. Maybe this is all above board, but I think it deserves a lot more scrutiny.

Subject Why are people fighting this?
Entry 03/20/2025 11:07 AM
Member RugPull1995

Look, I get and appreciate the desire to play devil's advocate and take the other side but I honestly don't get why people are fighting this one / have hope Coreweave will be a successful business. It isn't that hard:

- They have massive customer concentration with MSFT who is on record saying they only used CW at the beginning when AI became a thing since they couldn't react fast enough to build own DCs which have 12-24mo lead times. If anything Satya is signaling publicly several times that he is slowing down not speeding up on AI capex.
- Open AI said publicly stargate (led by ORCL) will satisfy 75%+ of OAI's compute needs through the end of the decade
- They are probably lying to you about GPU useful life - at the very least have to consider that AMZN just lowered useful life of GPUs to 5yrs from 6yrs and CW still at 6yrs and I would bet you that useful life keeps going down... Jensen literally told you at GTC that "once Blackwell is in full production, you won't be able to give away Hopper"
- They have a massive debt load with big principal payments that begin in 2026 through 2028
- They are relying on Core Scientific to build out a bunch of capacity for them, and Core Scientific has its own funding / resourcing problems
- The founders blew out of \$500m of stock prior to the IPO
- They got a deal with Open AI for "up to \$11B" in contract value in exchange for giving \$350m of stock. It is also unclear from my conversations with industry participants if that is all incremental or left pocket / right pocket with MSFT (meaning leaving MSFT b/s as commitment and going to OAI which is worse credit risk by a lot)
- They have locked in 3yr pricing for MSFT / other hyperscalers on GPUs and it is unclear if they can reset that pricing if tariffs cause servers / DC equipment to become more expensive

This thing isn't built to last. Their two largest customers are hyperscalers who compete with them and only gave them business since they didn't have enough capacity available.

The single most important metric for Coreweave in its quarterly earnings will be RPO - nobody will care about revenue since that is coming from RPO. They need to show significant RPO growth for the stock to work and I don't get how they do that given a huge portion of that RPO comes from large hyperscale customers who I think are unlikely to give CW future business given they are bringing their own capacity online.

Subject Re: Why are people fighting this?
Entry 03/20/2025 11:41 AM
Member gtwt

I side with the bears here. The main question on my mind though is, do I want to feel like a hero (hopefully successfully) shorting CRWV or keep betting against clear #2s and #3s in the space like WULF, EOSE, etc. who are reliably in an even worse position?..

Subject Re: Why are people fighting this?
Entry 03/20/2025 12:15 PM
Member Soprano14

I get the skepticism but I think folks here may be too negative on the OpenAI deal. In the roadshow video they seem to more clearly suggest this will add ~\$11bn in revenue commitments. Is there a reason to doubt this figure so significantly - has anyone heard anything more on if this is just MSFT's contract being taken over by OAI? thanks a lot

Subject Re: Re: Why are people fighting this?
Entry 03/20/2025 01:33 PM
Member lars

Well aside from the fact that the number might not be real / enforceable, and to RugPull1995's point it might be left pocket / right pocket, how do you get confidence that OpenAI can fund that? They are subject to the same external funding reflexivity that CRWV is. Doesn't seem like you're anywhere close to being compensated for the amount of uncertainty here.

Subject Re: Re: Why are people fighting this?
Entry 03/20/2025 03:39 PM
Member RugPull1995

I mean cmon!

<https://www.semafor.com/article/03/20/2025/microsoft-chose-not-to-exercise-12-billion-coreweave-option>

Subject Re: Re: Re: Why are people fighting this?
Entry 03/20/2025 04:20 PM

Member

tahoe88

Does anyone know how "contracted" all these RPOs are? Are there clauses that allow customers to break these? That seems like one of the more obvious situations where a pull back could cause a wework-esque collapse. If not, I imagine these guys will have a certain amount of runway and buys them time even if this is a horrible business model

Subject Re: Re: Re: Re: Why are people fighting this?
Entry 03/20/2025 09:10 PM
Member onodacapital

Contracts are take or pay, can exit for breach. Coreweave could of course totally fail to do their end, but they're legitimate contracts.

Again, I would not invest - and very well might short it. But the story the bulls will hear (again, the messenger) is 1) Nvidia wants this cloud to exist at scale and we may very well be structurally short compute for 5-10 years, and 2) they have expertise / differentiation at running AI-centric DCs. The numbers are also going to go up for 3-4 years and they will announce more deals. They will likely be doing >\$10B in revenue with 20% EBIT margins in 2-3 years (Cap structure & depreciation schedules are for another time). The bear case is that there's zero terminal value - probably true - but I've seen a lot dumber things in the last 5 years.

IPO apparently oversubscribed after day 1.

<https://www.bloomberg.com/news/articles/2025-03-21/coreweave-ipo-is-said-to-be-oversubscribed-after-first-day?srnd=homepage-americas>

Subject Re: Re: Re: Re: Re: Why are people fighting this?
Entry 03/20/2025 10:31 PM
Member lars

When has a banker not said an offering was oversubscribed, and yes, I'm sure Fidelity is in with a big order given they're already involved. Anyways, we've beat this one to death. Boggles my mind that people find an attractive R/R on this one. Appreciate the dialogue.

Subject Re: Re: Why are people fighting this?
Entry 03/24/2025 01:40 PM
Member Soprano14

Mgmt just confirmed the OpenAI contract is additive to the current RPO. Also said MSFT is all take-or-pay.

Subject Re: Re: Re: Why are people fighting this?
Entry 03/24/2025 06:43 PM
Member angus309

Soprano...Did you speak to the company and / or are your points somewhere in pr or public domain? Thank you.

Subject Re: Re: Re: Re: Why are people fighting this?
Entry 03/25/2025 08:01 AM
Member lars

<https://www.bloomberg.com/news/articles/2025-03-25/alibaba-s-tsai-warns-of-a-bubble-in-ai-datacenter-buildout?embedded-checkout=true>

Subject Re: Re: Re: Re: Re: Why are people fighting this?
Entry 03/26/2025 08:56 AM
Member lars

To RugPull1995's point on why force this - <https://www.ft.com/content/cb94eb68-ccb5-4fb3-b903-0aae17b836dd>

Not to mention: "Nvidia is CoreWeave's biggest supplier, one of its largest customers and an investor in the company."

Subject Re: Re: Re: Re: Re: Re: Why are people fighting this?
Entry 03/26/2025 11:21 AM
Member lars

Getting a lot of downvotes, so maybe none of this is appreciated I should quit posting (sorry!), but just seems like a lot of pertinent news flow...

<https://www.bloomberg.com/news/articles/2025-03-26/microsoft-abandons-more-data-center-projects-td-cowen-says?srnd=homepage-americas&embedded-checkout=true>

Subject Re: Re: Re: Re: Re: Re: Re: Why are people fighting this?
Entry 03/27/2025 08:27 AM
Member lars

<https://www.ft.com/content/a96d65c1-8ba7-4055-b952-61c5b5d500fc>

Subject Re: Re: Re: Re: Re: Re: Re: Re: Why are people fighting this?
Entry 03/27/2025 08:52 AM
Member gtwt

*COREWEAVE TO DOWNSIZE ITS IPO, CUT VALUE TO \$23B

Subject Re: Re: Re: Re: Re: Re: Re: Re: Re: Why are people fighting this?

Entry	03/27/2025 09:01 AM
Member	lars
Why wouldn't a downsized IPO cause a "run on the book" so to speak? They're so reliant on external funding, a smaller raise makes the equity considerably riskier (in my opinion, of course).	
Subject	Re: Re: Re: Re: Re: Re: Re: Re: Re: Re: Why are people fighting this?
Entry	03/27/2025 09:03 AM
Member	gtwt
Will be surprised if it goes through frankly at this point	
Subject	Re: Re: Re: Re: Re: Re: Re: Re: Re: Re: Why are people fighting this?
Entry	03/27/2025 11:45 AM
Member	gtwt
Seeing some contrarian views on CRWV on social media highlighting extreme investor scepticism, parallels with ARM's tight float and appeal to pods as reasons for a medium-term price upside. IMO ARM (whose stock is merely trading water itself) is an infinitely higher quality business than CRWV which seems like an ultimate zero. Skepticism in this case appears justified. Can CRWV trade up in the medium-term on technicals/"set up"/given Soprano's comment below on revenue visibility? Sure, but anyone with a 1 year+ view wouldn't care that much, and fundamentals >> "set up" big time in the long run.	
Subject	Re: Re: Re: Re: Re: Re: Re: Re: Re: Re: Why are people fighting this?
Entry	03/27/2025 11:53 AM
Member	lars
I'd counter with ARM's financial model being the polar opposite of CRWV - i.e. capital-light, very high return. Rather than using ARM to inform my view on CRWV, I would be more inclined to go the other direction and have my thoughts on CRWV inform my view on ARM - i.e. does someone like Softbank have to use a good asset like ARM in order to fund commitments / holes elsewhere. That seems to be a common dynamic in periods of financial distress.	
Subject	Re: Re: Re: Re: Re: Re: Re: Re: Re: Re: Why are people fighting this?
Entry	03/28/2025 09:58 AM
Member	gtwt
Not able to get any borrow today, this is unusual, a sign of either high short-seller demand or low institutional ownership - likely both.	
Subject	Re: Re: Re: Re: Re: Re: Re: Re: Re: Re: Why are people fighting this?
Entry	03/28/2025 10:45 AM
Member	RugPull1995
Top 3 accounts got half the deal (NVDA and two others who I am guessing are firm hands) and top 15 got 90% of the deal... so probably you are looking at 20m shares of real float or <\$1B so however this trades at open it won't be a real price	
Subject	Re: Re: Re: Re: Re: Re: Re: Re: Re: Re: Why are people fighting this?
Entry	03/28/2025 11:56 AM
Member	lars
Odds of the other two being Fidelity and Coatue? High, I'd say. Insiders / holders trying to frame this as an ARM-type scenario, but opposite environment and financial model. More likely a Venture Global. On that note, some more entertaining journalism from the FT - https://www.ft.com/content/11d04b02-081e-4df8-a7ff-22dc240c1ab4	
Subject	CRWV stock
Entry	04/01/2025 03:48 PM
Member	angus309
Is there any other logical explanation for the 30%+ move up today, outside of real institutional demand for CRWV stock? Genuinely asking, as this shouldn't be a squeeze per se (no borrow I thought), or can it? Asking what mechanics could explain this move outside of genuine demand. New eyes might be tempted to say GPU rental isn't circling the drain yet.	
Subject	Re: CRWV stock
Entry	04/01/2025 04:12 PM
Member	gtwt
Yeah no borrow still and a big move. Given the similarities with NMAX (new IPO, small float, big move), these look like deliberate squeeze targets. By whom? Retail (NMAX = "new GME"), algos, pods, I guess.	
And why would there be stable borrow in CRWV if daily trading volumes exceed float (ex-three anchors), implying few LT holders, I guess.	
Also, if these were genuine demand, last week people were welcome.	
Subject	Re: Re: CRWV stock
Entry	04/01/2025 04:18 PM
Member	RugPull1995
It's a 5% float stock (e.g., not a real price) and there is often some passive buying of new IPOs plus some lunatic pods probably thought this would be a good idea to short straight out of the gate	

Subject Re: Re: Re: CRWV stock
Entry 04/01/2025 04:22 PM
Member gtwt
 RugPull1995, no borrow still other than some chuncks for as much as 900%, I hear.

Subject Re: Re: Re: Re: CRWV stock
Entry 04/01/2025 04:24 PM
Member RugPull1995
 Yeah I have heard the same, don't think there is a point in shorting this in the near-term. Let the thing settle down a bit, let some more float appear, let them get a few bullish press releases out etc. (heard some rumors they might announce a GOOG contract win).

RIVN came out at a \$80B market cap or something and when it went to \$30B market cap you could still have made 60%+ on a short.

Subject Re: Re: Re: Re: Re: CRWV stock
Entry 04/01/2025 04:43 PM
Member blaueskobalt
 Look at the call volume today. Nearly 50% of the float.

Subject Re: Re: Re: Re: Re: Re: CRWV stock
Entry 04/01/2025 11:03 PM
Member sondasy
 And you had Cathie Wood lead the lemmings off the cliff... she's baaaack

Subject Re: Re: Re: Re: Re: Re: Re: CRWV stock
Entry 04/02/2025 04:00 PM
Member angus309
as per RugPull yesterday :

*14:33 **News Bot:** Google (GOOG) in advanced talks to rent Nvidia (NDA) AI servers from CoreWeave (CRWV), via The Information*

Subject ClickHouse - \$6 Billion Valuation
Entry 05/11/2025 10:31 AM
Member agentcooper2120
 From the thesis,

The business is likely to go through a funding round this year post the AWS contract at a valuation above \$4 billion, putting Nebius's stake at \$1+ billion.

<https://www.theinformation.com/articles/snowflake-challenger-clickhouse-targets-6-billion-valuation>

Good call. Assuming some dilution, maybe \$6/share in NBIS? Imagine if they're beating Snowflake for real-time workloads, they probably IPO in the next year or two right?

Subject questions
Entry 05/20/2025 05:28 PM
Member differentiatedfractal31415

Great call! A 'bitter lesson' indeed, if you follow Sutton. Thanks in advance for answering some questions I highlighted below. I am sure there should be bulls and bears debating the thesis again here, so hope this serves as a start.

- How should one think of IRR math for hopper vs blackwell, and could software/services help enough to make up for declines in commodity hardware rental prices? I am playing around with Rugpull's GPU Economics math posted in comments here, accounting now for management's new 4 year useful life estimate - thought it would be good to hear your take.
- How are you modeling capex requirements/MW for them to get from here to the targeted 1GW scale capacity over the next several years?
- What is the mid-term - 'a few years' answer on the call? 3 to 5 years? There is a wide range of interpretations of managerial adverbs..
- Of all the VC-like bets here, ClickHouse seems closer to real, if the \$6B estimate still holds, so accounting for that stake and reducing others to 0, for 400MW capacity NBIS EV \$/MW comes to \$14M/MW. Of course, Coreweave has its float-related and retail investor dynamic- it's trading at >\$125M/MW for its 420MW active power- when the dust settles, what is a reasonable \$/MW (industry/cycle average?) that should apply here in your NBIS analysis? (This goes back to the IRR/NPV math)
- Finally, came across this research blog from Semianalysis - any views on merits/limitations highlighted here for the name?<https://semianalysis.com/2025/03/26/the-gpu-cloud-clustermax-rating-system-how-to-rent-gpus/>

Subject Re: questions
Entry 06/24/2025 04:22 PM

Member

AAOI

Hey differentiatedfractal, big fan and thanks for the engagement. Been swamped, but we think the stock still hasn't seen its day in the sun. I'll address your questions and update the valuation as new news has come to light.

That in mind, NBIS is getting these GPUs at a steep discount, so run \$25-30K through that IRR model earlier on in this thread (lol); on top of that, since NBIS does the buildout and optimization of the stack, they don't need to pay Dell etc. to assemble racks, i.e., they absorb the additional margin that others like CRWV pay for thanks to its vertical integration. Even if that's not the case, regardless of what they sell it for, at this moment compute is tight and these guys are getting an objectively attractive return on capital.

The thing is, I think this all misses the forest for the trees though. As AI inference gets going, you just need to run GPUs, and the latest and greatest GPU is not really necessary. So fully depreciated GPUs can still make money. If cloud was CPUs, AI Cloud is GPUs. Years ago bear theses were also on the CPU depreciation and commoditization story, "the refresh rate on CPUs will crowd out earnings". Over time performance per dollar improved and the rest is history. And while CPUs are cheaper than GPUs, the steeper capex profile here for AI cloud is being met with just as steep revenue increases.

Also, it needs to be pointed out that Nebius is running a cloud-like usage model while CRWV is using an instance rental model. I.e. customer pays for renting a physical GPU through CRWV whereas a customer pays for usage time and access to a GPU via NBIS. So NBIS is able to optimize usage with possibly multiple customers on a gpu depending on when and how it's used.

Point being, the implications of all of this should be clear here, as NBIS can be much more capital efficient which should allow it to underprice their bare metal peers on a structural basis. And with the capital raised to date (\$1B convert), we think they need another \$2B to build out to 500MW. We think the incremental funds will come from a monetization of AvRide, perhaps my favorite of the wonderful little non core assets here but that's a conversation for another day (the food delivery robots are particularly interesting but I digress). Keep in mind as well that Yandex had a \$2B buyout option for Uber's 30% position in their self driving business in 2021. Why wouldn't Uber come back in on the same terms but with the business having *real* robotaxis in Austin later this year?

As for \$/MW, who knows. You could argue CRWV is a squeezey name and the implied \$150+/MW is way too high, but then once you net out NBIS other businesses (as we'll do below) the \$50+/MW is too low (IMO). Split the baby and you have massive upside even with the recent run. No need to dial in specificity when your margin of safety is as rock solid as it is with NBIS.

But long story short, once they're scaled to 500MW, we expect them to grab an anchor customer and then do a big raise to reach 1GW, which probably requires another \$5B or more. This will be most likely debt financed once they have a solid revenue stream. At any rate, the company thinks with the 500MW they'll have on tap by early next year they can reach \$5B in revenue with 20%-30% **EBIT** margins within 2 years. They are a little vague when they say medium term but think it's very much a scale the physical footprint this year, monetize it next, which gets you into 2027 (excluding another raise to 1GW). And they likely keep scaling as customers pull them.

So as for the Margin of Safety. ClickHouse is now valued at \$6.35B and people are catching on to the fact that it's really an incredible business. Specifically tuned to real time analytics and (shocker) great for AI workload analytics. Wouldn't be surprised if they IPO in another year or two, maybe above a \$10B valuation. But let's mark NBIS 'stake at \$1.75B.

Next is Toloka, which has ScaleAI as the best comp. Certainly people have noticed META buying out ScaleAI's team. In fact, because of that we've heard some people are moving to Toloka. Very similar story to ScaleAI, but a few years behind. Bezos buying maybe over 20% of Toloka (valuation undisclosed) for \$70M misses the point. One, check out Bezos VC track record, which is kind of a case in point, and two, keep in mind this is a strategic partner to scale Toloka through AWS. If Toloka is a few years behind ScaleAI, the company could grow and be valued materially higher. ScaleAI was valued at \$29B on \$2B of revenue. If Toloka is at \$75M in revenue today, a tripling isn't out of the question. Maybe Toloka is valued at \$1B in a year's time or so.

AvRide we discussed above, if they can sell a portion for \$2B, maybe it's worth \$4B total. Probably worth a lot more but that seems plenty for our purposes here.

Then if we take what they have in DC capacity today and if they were to sell the 100MW they have today at \$50k/MW (CRWV at something like \$150/MW), puts you at \$5B. Plus \$2.4B in cash, maybe \$200M for TripleTen, less \$1B in debt puts you at a liquidation value of \$44/share.

If I'm looking for a cheap bet on a fast growing industry which needs a lot of spending to get there, we think this is the one to bet on. CEO has 90% of his net worth in this, and everything they've done thus far has been smart. Execution is key but so far so good. Hope this helps!

Best,

AAOI

Subject Entry Member https://www.reuters.com/business/nebius-signs-174-billion-dollar-ai-infrastructure-deal-with-microsoft-2025-09-08/	MSFT NBIS Deal 09/08/2025 05:25 PM agentcooper2120
--	--

Nice, there's the anchor customer. 1GW soon? Great call. Bull case in the thesis playing out?

Subject Re: MSFT NBIS Deal
Entry 09/09/2025 10:20 AM
Member Shoe

Anyone have a guess at the economics of this deal?

Subject ORCL CEO comments are noteworthy
Entry 09/09/2025 05:50 PM
Member angus309

Oracle CEO Safra Catz:"As a bit of a preview, we expect Oracle Cloud Infrastructure revenue to grow 77% to \$18 billion this fiscal year—and then increase to \$32 billion, \$73 billion, \$114 billion, and \$144 billion over the subsequent four years."

Subject AAOI - NBIS
Entry 04/19/2026 08:31 AM
Member angus309

NBIS, has been a rocket ride at over 4x in just over a year. Congratulations. Between new contracting, and SOP...ClickHouse in particular (which could monetize), Toloka, Avride, and TripleTenClickHouse where are you coming out on valuation now? Thank you.

Subject Re: AAOI - NBIS
Entry 04/20/2026 01:35 PM
Member AAOI

Thanks Angus! We have \$425/share on 3.75GW assuming a linear extrapolation of the equity dilution they've done so far, which overestimates it as they'll be more cash funded as the flywheels spin up. FWIW they think they can hit 5 GW in the next few years. Not sure I'd ever want to bet against Arkady, the guys track record of execution over decades is truly bonkers. In any case, we think we're being fairly conservative, and probably makes sense for us to state for the record that we think the first AI native hyper scaler will be Nebius. You can forget about Iren or CoreWeave and other guys - the one that stands apart in all this is clearly Nebius.

Subject Re: Re: AAOI - NBIS
Entry 04/20/2026 01:47 PM
Member angus309

Really appreciate it, AAOI. Great outcome so far, with more to come. Arkady's story is truly amazing...looks like roughly 20% short, which is a big mistake IMO. Arkady will find a way to win.

Subject Re: Re: Re: AAOI - NBIS
Entry 04/20/2026 01:50 PM
Member angus309

Forgot to add that ClickHouse should be interesting this year. Does your valuation target account for ClickHouse (and other investments)?

Subject vineland
Entry 04/24/2026 06:30 PM
Member Lyncroft

Any concerns on their nerfing the 2026 guide because of Vineland delays?